K2SN-MSS: An Efficient Post-Quantum Signature*

Sabyasachi Karati[†]

School of Computer Sciences

National Institute of Science Education and Research, India

skarati@niser.ac.in

Reihaneh Safavi-Naini

Department of Computer Science

University of Calgary, Canada

rei@ucalgary.ca

ABSTRACT

With the rapid development of quantum technologies, quantum-safe cryptography has found significant attention. Hash-based signature schemes have been in particular of interest because of (i) the importance of digital signature as the main source of trust on the Internet, (ii) the fact that the security of these signatures relies on existence of one-way functions, which is the minimal assumption for signature schemes, and (iii) they can be efficiently implemented. Basic hash-based signatures are for a single message, but have been extended for signing multiple messages. In this paper we design a Multi-message Signature Scheme (MSS) based on an existing One-Time Signature (OTS) that we refer to as KSN-OTS. KSN uses SWIFFT, an additive homomorphic lattice-based hash function family with provable one-wayness property, as the one-way-function and achieves a short signature. We prove security of our proposed signature scheme in a new strengthened security model (multi-target multi-function) of MSS, determine the system parameters for 512 bit classical (256 bit quantum) security, and compare parameter sizes of our scheme against XMSS, a widely studied hash based MSS that has been a candidate for NIST standardization of post-quantum signature scheme. We give an efficient implementation of our scheme using Intel SIMD (Single Instruction Multiple Data) instruction set. For this, we first implement SWIFFT computation using a SIMD parallelization of Number Theoretic Transform (NTT) of elements of the ring $\mathbb{Z}_p[X]/(X^{\hat{n}} + 1)$, that can support different levels of parallelization. We compare efficiency of this implementation with a comparable (security level) implementation of XMSS and show its superior performance on a number of efficiency parameters.

CCS CONCEPTS

• **Security and privacy → Digital signatures**.

KEYWORDS

OTS, Merkle Tree, NTT, SWIFFT, Cover-Free Family, SIMD.

1 INTRODUCTION

Cryptographic primitive Digital signatures [14] form the basis of establishing trust over the Internet. Today's commonly used signatures are RSA, DSA and ECDSA [12], and use the hardness of Integer Factorization and Discrete Logarithm problems as the basis of security. Shor [39] gave efficient quantum algorithms for both the problems which make these signature schemes insecure against quantum computers. The recent call by security agencies [40] has resulted in efforts by IETF [1], NIST [36] and other organizations to move to post-quantum algorithms, and this has made development of post-quantum secure digital signature of high importance. As the digital signatures form the backbone of trust on the Internet, it is essential to employ quantum-safe signature algorithms to extend this trust into the future. To provide security against quantum computers, digital signature schemes may use computational assumptions such as the hardness of decoding of random codes [9], or finding short vectors in lattices [15], for which no efficient quantum algorithm is known today. Such schemes will remain secure as long as no efficient quantum algorithms for the underlying problems are found. However, these signatures suffer from inefficient computation, large signature length and/or public keys, or unproven security [5].

Quantum-safe digital signature schemes can be based on One-Way-Functions (OWF) which are known to be the minimum assumption for digital signatures [38], and is expected to provide long-term security assurance. Hash-based signature schemes use hash function families as OWF. OTS, introduced by Lamport in 1979 [27], uses a secret and public key-pair to sign a single message. In Lamport's construction, the secret key is a set of random binary strings, called component secret keys, and the public key is obtained by applying a OWF to each component secret key. A message is associated with a subset of component keys and signature will be the set of secret component keys associated with the subset. Security of the scheme depends on the hardness of inverting the OWF. The main advantages of OTS are: (i) it can be built using any OWF, and (ii) the signing and the verification are fast [42]. The sizes of the keys and the signature, however, can be significantly larger than today's digital signatures such as RSA, DSA or ECDSA.

To sign 2^h messages, one needs to use a OTS scheme with 2^h different keys. Using a Merkle tree [31], the public keys of the 2^h OTSs can be hashed into a single short public key. A signature now will include the signature that is generated by a OTS, together with

a list of nodes in the Merkle tree which are required to compute the OTS public key to the root of the Merkle tree. A signed message will have an index/address that will indicate the position of OTS public key in the set of leaves of the Merkle tree, that will be used for authentication path verification.

In [11], a generic construction of a multi-message signature from a OTS using Merkle tree is given and the security of MSS is reduced to the security of the OTS and the second pre-image resistance of the hash function. The construction, known as SPR-MSS, is particularly attractive because it relaxes the requirement on the hash function (second pre-image instead of one-wayness and collision-freeness), and has been the basis of widely studied construction of multi-message signature scheme XMSS [7, 19]. XMSS (eXtended Merkle Signature Scheme) uses W-OTS$^+$ [32] as the one-time signature. XMSS is a prime candidate of standardization [35] of quantum-safe secure signature schemes by NIST [36] and IETF [1]. In [21], it was shown that the security model of MSS must be strengthened to realistically capture all the hash values that become available to the adversary in a MSS setting and can significantly improve their success chance of a forgery. The new security notion is captured by multi-target-multi-function preimage (or second-preimage) property in [21].

In [23], Kalach and Safavi-Naini introduced a hash-based OTS scheme, called as KSN-OTS, that reduces the number of secret key components of Lamport's system [27] by using a 1-CFF (Cover-Free-Family - See Definition 3.1). KSN-OTS has a short signature because of using a family of homomorphic (additive) hash functions called SWIFFT that allows the component secret keys of a message be combined. KSN-OTS security proof reduces unforgeability of the signature to the collision resistance of SWIFFT. KSN-OTS has short signature length, and fast signing and verification (Table 1 of [23]). The main computation step in KSN-OTS is the computation of SWIFFT function that needs $\hat{m}$ multiplications over $\mathcal{R} = \mathbb{Z}_p[X]/(X^{\hat{n}} + 1)$. The fastest method of implementing ring multiplication uses Number Theoretic Transform (NTT) [28] to transform the ring elements to vectors of dimension $\hat{n}$, and allows ring multiplication and addition to become component-wise vector multiplication and addition modulo p.

1.1 Our Contribution

We design K2SN-MSS, a multi-message signature scheme, using an approach that is inspired by SPR-MSS [11], and prove its security using the notion of multi-target, multi-function pre-image and second pre-image resistance. We use SWIFFT as the hash function for KSN-OTS and also the Merkle tree. This has the benefit of using a single optimized code for the whole construction. However it requires addressing a number of challenges, including the need to introduce a new operation, called Merge, for constructing the Merkle tree to compensate for the mis-match between parameter sizes of SWIFFT and the SPR-MSS construction. We prove that use of this operation does not affect security of the construction.

We implement the signature scheme using an efficient implementation of SWIFFT which relies on an efficient implementation of NTT. We use SIMD parallelization to achieve low level instructional efficiency. SWIFFT and KSN-OTS are highly parallelizable

algorithms that can be further parallelized at higher levels (process level).

To compare our results with XMSS, we will obtain concrete security parameters of K2SN-MSS taking into account security level and parameters of existing implementations of XMSS. Our results show that K2SN-MSS has 3 times faster key generation, signature generation and verification. It however has 5% larger public key and 4 times longer signature.

1.1.1 K2SN-MSS. Our construction uses a modified construction of SPR-MSS in two ways: (i) the Merge operation is used in each tree node to combine the hash output of the lower level, and (ii) different hash keys and random pads are used at each node and so each node is effectively using a different hash function. We prove that with these modifications, the resulting MSS is secure in the multi-function multi-target attack model, and the security of the scheme reduces to the multi-function multi-target preimage and second-preimage resistance of the underlying hash function.

1.1.2 K2SN-MSS implementation. K2SN-MSS consists of three algorithms: key generation, signing and verification algorithms. To generate component secret keys, hash keys and random pads of the Merkle tree, ChaCha20 pseudo-random function family (PRF) is used. ChaCha20 is a state-of-the-art PRF which takes 40-bytes input seed and can generate output of 128 bytes long [4]. We choose parameters of SWIFFT such that it provides 512-bit classical (256-bit quantum) security for K2SN-MSS against existential unforgeability in chosen message attack (EUF-CMA) (See Section 3.3). We implement the following code modules

- Key generation algorithm, that uses Chacha20 as a submodule, and computes the component secret keys, hash keys and the random pads. SWIFFT hash function was used to compute the component public keys, and construct the Merkle tree.
- We implemented the 1-CFF Algorithm [6] to determine the subset of component keys that are associated with a message. This module is called in signing and verification, both.
- The signing and the verification algorithms use ChaCha20, SWIFFT, and the 1-CFF algorithm above.

The time for key generation, signing and verification is dominated by the time required for calculating the hash function SWIFFT. Below we briefly describe our efficient implementation of SWIFFT using SIMD instructions.

1.1.3 Efficient implementation of SWIFFT using SIMD. Ring multiplication over the ring $\mathcal{R}$ dominates the computation cost of SWIFFT function. For the efficient implementation of multiplication in the ring $\mathcal{R}$, we use NTT. We show that the computation of NTT transformation can be parallelized using SIMD, with different levels of parallelization that is parametrized by δ, denoting simultaneous computation of 2^δ components of the output vector. The parallelization can be used for $\delta = 0$ (no parallelization) to $\delta = 6$ (full parallelization for the ring $\mathcal{R}$ with $\hat{n} = 64$). The parallelization in [2, 30, 34] is an instance of our general approach with $\delta = 3$. We show the result of SWIFFT function with $\delta = 4$ that leads to 25% higher speed compared to [34]. Our implementations of NTT computation can be with, or without, precomputation. The former approach results in improved speed but the memory requirement

grows double-exponentially $O(2^{2^\delta})$ in the value of δ. Thus we consider precomputation when the required memory is acceptable.

For efficient verification of KSN-OTS (and also K2SN-MSS), we also implemented gSWIFFT (generalized input SWIFFT) that allows non-binary input strings. This extension was introduced in [23] and is used in our work.

1.1.4 Evaluation of K2SN-MSS. We compare performance of K2SN-MSS with XMSS for 2^h messages, each of length 256 bits, and security parameter $n = 512$-bit. We use SWIFFT function with parameters $\hat{m} = 16$, $\hat{n} = 64$, $p = 257$ (Sec Section 2.4). As we will discuss in Section 2.1, these parameters allow us to have a fair comparison in terms of functionality and security level, with XMSS. Additionally, they allow us to compare our implementation of SWIFFT with the existing SWIFFT implementation in [30].

It shows that K2SN-MSS has smaller processing time for all algorithms: key generation, signing and verification. The corresponding times for K2SN-MSS are 2.76, 2.89 and 2.65 times faster than XMSS. On the other hand, the signature length of K2SN-MSS is approximately 4 times larger than that of XMSS. Public and secret key sizes of the two schemes are comparable (almost same). Thus the two schemes will have complementary speed/signature size properties, while enjoying similar security guarantee.

Table 1: Comparison between XMSS and K2SN-MSS scheme at 512-bit classical security level

Name	K2SN-MSS/SWIFFT-16-avx2	XMSS/SHA512/$w = 16$ [13]
Key Generation (ms)	526069	1452600
Signing (ms)	4.70	13.57
Verification (ms)	0.34	0.90
Signature Size	21331 Bytes	5571 Bytes
Secret Key Size	Seed	Seed
Public Key Size	152 Bytes	144 Bytes

We note that although we implemented KSN-OTS using a single hash function, but the homomorphic property of SWIFFT is only required in the verification of KSN signatures and one can use a different hash function, such as SHA2-512, for the construction of the Merkle tree. Our choice of SWIFFT for Merkle tree is further discussed in Section 7.

2 PRELIMINARIES

We use the following notations and function definitions.

- $n \in \mathbb{N}$: Security parameter where $\mathbb{N}$ denotes the set of natural numbers.
- $m \in \mathbb{N}$: Length of the messages in bits.
- $\hat{m}, \hat{n} \in \mathbb{N}$: Parameters of SWIFFT hash function family.
- p: Prime integer.
- t: The number of component key-pairs for the OTS we employ. We assume $t = 2^\ell$ for some positive integer ℓ and $\log_2\left(\frac{t}{2}\right) \geqslant m$.
- h: Merkle tree height. The K2SN-MSS is designed to sign 2^h messages.
- $\|$: concatenation of two binary strings.
- $\oplus$: bit-wise XOR of two binary strings.

- $\oplus_p$: component-wise modulo p addition of two vectors of same length.
- $\odot_p$: component-wise modulo p multiplication of two vectors of same length.
- $F_n = \left\{ f_k : \{0,1\}^{2n} \to \{0,1\}^{2n} \mid k \in \{0,1\}^n \right\}$ is pseudo-random function family and will be used in KSN-OTS and K2SN-MSS.
- $\mathcal{F}_n = \{ f_k : \{0,1\}^n \to \{0,1\}^n \mid k \in \{0,1\}^n \}$ is pseudo-random function family and will be used in W-OTS$^+$ and XMSS.
- c_m: A function that, on an input of m-bit, generates a row of the 1-CFF function family table.

2.1 XMSS and W-OTS$^+$

XMSS [7] is a hash-based signature scheme that extends the one time signature W-OTS$^+$ to a multi-message signature scheme using Merkle tree construction. The construction was proposed in [7] and the security was reduced to the second-preimage resistance of the hash function that was used in the Merkle tree and the security of the W-OTS$^+$. XMSS^{MT} [20] extends XMSS to multiple layer tree where each layer root node is signed using W-OTS$^+$. It *allows for a virtually unlimited number of signatures, while it preserves its desirable properties* [20] of XMSS. This provides signature and key generation computation time and signature size trade-off which allow the users to choose parameters to match their requirements. XMSS is a special case of XMSS^{MT}. In [21], XMSS is proposed that provides security against multi-function-multi-target preimage (or second preimage) attacks. The new security requirement for the signature schemes is given in full version [25] and is used in our security analysis. There are several implementations for XMSS [13, 18]. We compare our implementation with the implementation in [13] because this takes the multi-function-multi-attack notion of security into consideration. Also, this is the only implementation which takes advantage of SIMD parallelization and reports the fastest result. On the other hand, choice of the parameters for SWIFFT hash function leads to 512-bit classical (256-bit quantum) security of K2SN-MSS in multi-function-multi-target model and it should be noted that there is no existing software of XMSS^{MT} at the same security level.

2.2 W-OTS$^+$

XMSS signature scheme uses W-OTS$^+$ [7] as the OTS scheme. The public key and signature of W-OTS$^+$ is computed through random walks in a function family $\mathcal{F}_n = \{ f_k : \{0,1\}^n \mapsto \{0,1\}^n \mid k \in \{0,1\}^n \}$. For a $k, x \in \{0,1\}^n$ and $e \in \mathbb{N}$, the random walk is defined as:

$$f_k^0(\mathbf{x}) = k, \quad f_k^i(\mathbf{x}) = f_{k'}(\mathbf{x}), \quad \text{where } k' = f_k^{i-1}(\mathbf{x}).$$

For an m-bit message, and *Winternitz parameter w*, define following parameters

$$l = l_1 + l_2, l_1 = \left\lceil \frac{m}{\log_2(w)} \right\rceil, l_2 = \left\lfloor \frac{\log_2(l_1(w-1))}{\log_2(w)} \right\rfloor + 1. \quad (1)$$

The secret key of W-OTS$^+$ is $\mathcal{SK} = \{\mathbf{x}_0, \mathbf{x}_1, \ldots, \mathbf{x}_{l-1}\}$ where each $\mathbf{x}_i$ is a binary string of length n (that is generated by using a pseudo-random function on a random seed), and the corresponding public key is $\mathcal{PK} = \{\mathbf{x}, f_{\mathbf{x}_0}^{w-1}(\mathbf{x}), f_{\mathbf{x}_1}^{w-1}(\mathbf{x}), \ldots, f_{\mathbf{x}_{l-1}}^{w-1}(\mathbf{x})\}$, where $\mathbf{x}$ is a randomly generated element from $\{0,1\}^n$.

To sign a message mes, first the message is written in base w as, mes $= \{b_0, b_1, \ldots, b_{l_1-1}\}$. Next a checksum $C = \sum_{i=0}^{l_1-1}(w - 1 - b_i)$ in base w is calculated, $C = \{b_{l_1}, b_{l_1+1}, \ldots, b_{l-1}\}$, and appended to the message to obtain $b = \{b_0, b_1, \ldots, b_{l-1}\}$.

The signature of b is $\sigma_{\text{mes}} = \{\sigma_0, \sigma_1, \ldots, \sigma_{l-1}\} = \{f_{\mathbf{x}_0}^{b_0}(\mathbf{x}), f_{\mathbf{x}_1}^{b_1}(\mathbf{x}), \ldots, f_{\mathbf{x}_{l-1}}^{b_{l-1}}(\mathbf{x})\}$, resulting in an $l \times n$-bit signature. Verification is done by checking

$$\mathcal{PK} \stackrel{?}{=} \{f_{\sigma_0}^{w-1-b_0}(\mathbf{x}), f_{\sigma_1}^{w-1-b_1}(\mathbf{x}), \ldots, f_{\sigma_{l-1}}^{w-1-b_{l-1}}(\mathbf{x})\}.$$

2.3 XMSS: A Secure Merkle Tree MSS Construction [21]

In this section we describe XMSS briefly which is an instantiation of SPR-MSS with W-OTS$^+$ and also takes the multi-function-multi-target security model of hash functions into consideration. Let W-OTS$^+$ be a secure one-time signature scheme with 2^ℓ component secret keys (and the corresponding component public keys). To design an MSS for 2^h messages, 2^h instances of the W-OTS$^+$ will be used as follows. Let the secret key $\mathcal{SK}_i$ and public key $\mathcal{PK}_i$ of the i-th instance of the W-OTS$^+$ be,

$$\mathcal{SK}_i = \{\mathbf{x}_{i_0}, \mathbf{x}_{i_1}, \ldots, \mathbf{x}_{i_{2^\ell-1}}\}; \ \mathbf{x}_{i_j} \in \{0,1\}^n, \text{ and}$$
$$\mathcal{PK}_i = \{\mathbf{x}_i, \mathbf{y}_{i_0}, \mathbf{y}_{i_1}, \ldots, \mathbf{y}_{i_{2^\ell-1}}\}; \mathbf{x}_i, \mathbf{y}_{i_j} \in \{0,1\}^n.$$

To sign the i-th message, the i-th instance of the W-OTS$^+$ is used. The construction uses a Merkle tree of height $h + \ell$, and attaches a random binary string of length $2n$ for each node of the tree, denoted by $\mathbf{v}_{i,j}$ where (i, j) denotes the j-th node from left at height i as $0 \leqslant i \leqslant h + \ell$ and $0 \leqslant j \leqslant 2^i - 1$. In our construction, we consider that leaf nodes are at height $h + \ell$ and the root node is at the height 0. Similarly each node computation uses a different hash key $k_{i,j}$ and random pad $\mathbf{v}_{i,j}$. The construction of the Merkle tree consists of two types of trees: (i) $\mathcal{L}_i$ trees and (ii) $\mathcal{MSS}$ tree.

(1) Construction of $\mathcal{L}_i$ trees: $\mathcal{L}_i$ is a Merkle tree of height ℓ whose leaves are the component public keys of the i-th instance of W-OTS$^+$, that is $\{\mathbf{x}_i, \mathbf{y}_{i_0}, \mathbf{y}_{i_1}, \ldots, \mathbf{y}_{i_{2^\ell-1}}\}$. Let $\mathbf{y}_{i_{l,j}}$ be a node of the tree where it is the j-th ($0 \leqslant j \leqslant 2^l - 1$) node from the left at height l ($0 \leqslant l \leqslant \ell$) of the $\mathcal{L}_i$. Now we assign $\mathbf{y}_{i_{\ell,j}} = \mathbf{y}_{i_j}$ for $0 \leqslant j < 2^\ell$. The nodes $\mathbf{y}_{i_{l,j}}$ for $0 \leqslant l < \ell$ are computed as

$$\mathbf{y}_{i_{l,j}} = H_{k_{i_{l,j}}}((\mathbf{y}_{i_{l+1,2j}} \| \mathbf{y}_{i_{l+1,2j+1}}) \oplus \mathbf{v}_{i_{l,j}}), \tag{2}$$

where $H_k \in \mathcal{H}_n$, $\|$ and $\oplus$ denote the concatenation and bit-wise XOR of two binary strings, respectively. Let the root of the Merkle tree $\mathcal{L}_i$ be $\mathbf{y}_{h,i}$ and then $\mathbf{y}_{h,i} = \mathbf{y}_{i_{0,0}}$.

(2) Construction of $\mathcal{MSS}$ tree: $\mathcal{MSS}$ has height h and the leaf nodes are the root nodes of the $\mathcal{L}_i$ trees, that is $\{\mathbf{y}_{h,0}, \mathbf{y}_{h,1}, \ldots, \mathbf{y}_{h,2^h-1}\}$. The intermediate nodes and the root node of the $\mathcal{MSS}$ are constructed in the similar way shown in equation (2). Let the root of the $\mathcal{MSS}$ tree be $\mathbf{y}_{0,0}$.

The pictorial view of the Merkle tree is given in the Figure 1 and Figure 2.

All the random pads and all the hash keys can be generated using a seed and pseudo-random number generator (PRG). From the seed and the PRG, we can compute the random pad and the hash key for a particular node using the index of the node in the


Figure 1: Merkle Tree

Figure 2: Intermediate node computation

tree. Let the seed for hash key be hkseed and the seed of random pads be rpseed. Also all the secret keys can be generated using the PRG using another seed sk. The multi-message signature scheme has secret and public keys given by, $\mathcal{SK} = \{\text{sk}\}$ and $\mathcal{PK} = \{\mathbf{y}_{0,0}, \text{hkseed}, \text{rpseed}\}$, respectively. A signature consists of, (i) the index of the message, i, (ii) the signature of the i-th W-OTS$^+$ on the message, (iii) the public key $\mathcal{PK}_i$ of the i-th W-OTS$^+$, and (iv) a list of nodes of $\mathcal{MSS}$ called *authentication path* which authenticates the root of the $\mathcal{L}$ tree that corresponds to $\mathcal{PK}_i$, against the root of the $\mathcal{MSS}$ tree, $\mathbf{y}_{0,0}$. As the public key of W-OTS$^+$ can be computed using the random walks in the function family $\mathcal{F}_n$ from the W-OTS$^+$ signature, the XMSS signature does not include the aforementioned third component $\mathcal{PK}_i$ of i-th instance of W-OTS$^+$. K2SN-MSS uses this structure of Merkle, and employ a pseudo-random generator together with a random seed to generate the secret key components. Hash keys and random pads are also generated in a similar way.

2.4 SWIFFT: An Efficient Lattice-Based Hash-Function

The SWIFFT family of hash functions [30] is defined by three parameters: (i) $\hat{n}$: a power of 2, (ii) $\hat{m}$: a small integer, and (iii) p: a prime (the construction works for non-prime also).

Let $\mathcal{R}$ be a ring as defined in the equation (3):

$$\mathcal{R} = \mathbb{Z}_p[X]/(X^{\hat{n}} + 1). \tag{3}$$

A SWIFFT function of a given SWIFFT family is defined by $\hat{m}$ elements, $\mathbf{a}_0, \mathbf{a}_1, \ldots, \mathbf{a}_{\hat{m}-1} \in \mathcal{R}$ called *multipliers*. The input of a SWIFFT function is a binary string of length $\hat{m}\hat{n}$: the input is divided into $\hat{m}$ $\hat{n}$-bit string, and each string is represented as an element of $\mathcal{R}$ with binary coefficients. Let $\mathbf{z} = (\mathbf{z}_0, \mathbf{z}_1, \ldots, \mathbf{z}_{\hat{m}-1})$ be a SWIFFT input, where $\mathbf{z}_i$ is a ring element with binary coefficients. Then

hash of $\mathbf{z}$ is computed as,

$$\text{SWIFFT}(\mathbf{z}) = \sum_{i=0}^{\hat{m}-1} (\mathbf{a}_i \cdot \mathbf{z}_i), \tag{4}$$

where the multiplications and the additions are in the ring $\mathcal{R}$. For computational efficiency, the ring multiplications are performed using NTT and so SWIFFT is written as equation (5):

$$\text{SWIFFT}(\mathbf{z}) = \sum_{i=0}^{\hat{m}-1} \left(\text{NTT}(\mathbf{a}_i) \odot_p \text{NTT}(\mathbf{z}_i) \right). \tag{5}$$

Notice that we do not apply iNTT on the output of equation 5 to get back the same result as of equation 4. As FFT is linear and bijective, there is no need to apply iNTT on the output of equation 5 as suggested in [30].

Security. SWIFFT is an efficient instantiation of the generalized compact knapsack (GCK) [33] [30]. [33] proves that the GCK is asymptotically one-way, and this implies that the SWIFFT family also enjoys the same property. GCK is proved to be asymptotically collision resistant[29, 37]. In [30], it is proved that the security of a SWIFFT function reduces to the security of the subset-sum problem.

2.4.1 gSWIFFT. We use a generalized version of SWIFFT function that takes input from $\mathbb{Z}_p^{\hat{m}\hat{n}}$ (instead of $\{0, 1\}^{\hat{m}\hat{n}}$). Let the input to gSWIFFT be $\mathbf{z}' = (\mathbf{z}'_0, \mathbf{z}'_1, \ldots, \mathbf{z}'_{\hat{m}-1})$ where $\mathbf{z}'_i \in \mathbb{Z}_p^{\hat{n}}$. The input length of gSWIFFT is $\hat{m}\hat{n} \lceil \log_2(p) \rceil$ bits. Similar to SWIFFT, we define gSWIFFT function as:

$$g\text{SWIFFT}(\mathbf{z}') = \sum_{i=0}^{\hat{m}-1} (\text{NTT}(\mathbf{a}_i) \odot_p \text{NTT}(\mathbf{z}'_i)), \tag{6}$$

gSWIFFT is used for efficient verification of KSN-OTS scheme (Section 3.1).

2.4.2 Parallelizing SWIFFT Implementation using SIMD. Efficient implementation of K2SN-MSS requires efficient implementation of SWIFFT computation that consists of $\hat{m}$ ring multiplications, followed by ring additions. Thus efficient implementation of ring multiplication is essential for the efficient implementation of K2SN-MSS. We show how to parallelize each ring multiplication using SIMD instructions of Intel Intrinsics. The level of parallelization is determined by the parameter δ, which indicates 2^δ output components will be computed simultaneously. In our setting, the parallelization level δ, can be chosen in the range $\delta = 0$ (no parallelization at all) to $\delta = 6$ corresponding to the full parallelization for ring $\mathcal{R}$ with $\hat{n} = 64$.

Intel Intrinsics for SIMD Intel introduced SIMD instruction set MultiMedia eXtension (MMX) in their processors in 1997. MMX has been later extended to Streaming SIMD Extensions (sse) like sse2, sse3 upto sse4.2. We use Advanced Vector Extensions (avx) and in particular avx2 intrinsic. For more details see [16]. The achievable level of parallelization will depend on the specific set of instructions that will be chosen for the implementation. For 16 fold parallelization ($\delta = 4$), we use 16 16-bit multiplication SIMD instruction together with 16-bit addition, subtraction and other operations such as shift. More details are in Section 5.

3 K2SN-MSS CONSTRUCTION

We first give an overview of KSN-OTS [23] and then describe K2SN-MSS, and prove its security in using the stronger notion of multi-target multi-target security.

3.1 KSN: An Efficient Post-Quantum OTS

We propose an MSS based on KSN [23], an OTS scheme with two building blocks: a cover-free-family and SWIFFT hash function family.

A *w-Cover-Free Family (w-CFF)* is a set system defined as follows.

Definition 3.1. Let $(X, \mathcal{B})$ be a set system, where $X = \{x_1, x_2, \ldots, x_t\}$ is a set of t elements and $\mathcal{B} = \{B_1, B_2, \ldots, B_e\}$ be a set of e subsets of X and the size of each B_i is k. If for all $\Delta \subset \{1, 2, \ldots, e\}$ with $|\Delta| = w$ and for all $i \notin \Delta$, we have $\left| B_i \smallsetminus \bigcup_{j \in \Delta} B_j \right| \geqslant 1$ then we call that the set system is a k-uniform w-Cover-free family with t elements and e subsets, in short w-CFF(t, e, k).

The SWIFFT hash function family is additive homomorphic and for two binary inputs $\mathbf{x}_1$ and $\mathbf{x}_2$ of length $\hat{m}\hat{n}$, satisfying

$$\text{SWIFFT}(\mathbf{x}_1) + \text{SWIFFT}(\mathbf{x}_2) = g\text{SWIFFT}(\mathbf{x}_1 \oplus_p \mathbf{x}_2),$$

where addition $\text{SWIFFT}(\mathbf{x}_1) + \text{SWIFFT}(\mathbf{x}_2)$ is in ring $\mathcal{R}$, and addition $\mathbf{x}_1 \oplus_p \mathbf{x}_2$ is a component-wise modulo p addition of two vectors of length $\hat{m}\hat{n}$. This property is used in KSN signature scheme [23] to achieve fast signing and verification, and short signature. An overview of the scheme is given below. We use a pseudo-random function family $F_n : \{f_k : \{0, 1\}^{2n} \mapsto \{0, 1\}^{2n} \mid k \in \{0, 1\}^n\}$ to generate component secret keys (instead of uniform selection from $\{0, 1\}^{2n}$).

Consider a message space consisting of m-bit strings. The secret key and the public key of KSN-OTS scheme consists of t component keys, where t is an even integer satisfying $\log_2 \left(\frac{t}{2} \right) \geqslant m$. A component secret is a $2n$ bit random string and so the total secret key length is $2tn$ bits. The component secret keys in practice (and in our implementation) are generated by applying $f \in F_n$ on an initial seed. Thus the secret key can be considered as the seed of the PRG. The hash key k of SWIFFT is $\hat{m}$ multipliers which are ring elements and are chosen uniformly at random from $\{0, 1\}^{\hat{n} \lceil \log_2(p) \rceil}$.

The signature scheme consists of three algorithms: KeyGen, Sign and Verify. The parameters of SWIFFT function ($\hat{m}$, $\hat{n}$ and p) and the pseudo-random function family F_n will be chosen to achieve the required security level. A function $f_{sk}()$ is chosen from F_n, and sk is kept secret (this can be seen as the seed of the pseudo-random function).

Key Generation. The KeyGen generates the secret key $\mathcal{SK}$ and the public key $\mathcal{PK}$ as given in Algorithm 1. The i-th component secret key is computed as $f_{sk}(i + 1)$. This is the same as the secret key generation in XMSS [7].

Algorithm 1 Key Generation (KeyGen(1^n))

---

(1) Choose sk as a seed uniformly at random from $\{0, 1\}^n$.
(2) For $i = 0, 1, \ldots, t - 1$, compute $\mathbf{x}_i$ as $\mathbf{x}_i = f_{sk}(i + 1)$.
(3) For $i = 0, 1, \ldots, t - 1$, compute $\mathbf{y}_i$ as $\mathbf{y}_i = \text{SWIFFT}_k(\mathbf{x}_i)$.
(4) Return Secret Key $\mathcal{SK}$ = sk and Public Key $\mathcal{PK}$ = $\{\mathbf{y}_0, \mathbf{y}_1, \ldots, \mathbf{y}_{t-1}\}$.

Signature Generation. The Sign takes an m-bit message mes and computes the signature σ_{mes} using Algorithm 2.

Signature Verification. The verification algorithm given by Algorithm 3, takes a message and signature pair (mes, σ_{mes}), and outputs Accept or Reject.

Algorithm 2 Signature Generation ($\texttt{Sign}(\text{mes}, \mathcal{SK})$)

(1) For $i = 0, 1, \ldots, t - 1$, compute $\mathbf{x}_i$ as $\mathbf{x}_i = f_{\mathcal{SK}}(i + 1)$.
(2) Compute $B_{\text{mes}} = c_m(\text{mes})$, where $B_{\text{mes}} = \{i_0, i_1, \ldots, i_{\frac{t}{2}-1}\}$, and $i_j < i_{j+1}$ for all $0 \leqslant j \leqslant \frac{t}{2} - 2$.
(3) Compute $\sigma_{\text{mes}} = \mathbf{x}_{i_0} \oplus_p \mathbf{x}_{i_1} \oplus_p \cdots \oplus_p \mathbf{x}_{i_{\frac{t}{2}-1}}$; Return σ_{mes}.

Algorithm 3 Signature Verification ($\texttt{Verify}(\text{mes}, \sigma_{\text{mes}}, \mathcal{PK})$)

(1) Compute $B'_{\text{mes}} = c_m(\text{mes})$.
(2) If $\sigma_{\text{mes}}[i] \leqslant \frac{t}{2}, \forall 0 \leqslant i < \hat{m}\hat{n}$ ($\sigma_{\text{mes}}[i]$ denotes the i-th component of the vector σ_{mes}), then continue else Reject.
(3) Compute $\sigma'_{\text{mes}} = \mathbf{y}_{i_0} + \mathbf{y}_{i_1} + \cdots + \mathbf{y}_{i_{\frac{t}{2}-1}}$ (addition in $\mathcal{R}$), and $i_j \in B'_{\text{mes}}$.
(4) If $\sigma'_{\text{mes}} = g\text{SWIFFT}_k(\sigma_{\text{mes}})$, then Accept, else Reject.

In [23], it has been proved that the KSN scheme is secure under the collision-resistance of the hash function. The security theorem is as given below:

Theorem 3.2. *[23] If $\mathcal{A}$ is a quantum adversary that (ϵ, τ)-wins the unforgeability security game for KSN scheme, then $\mathcal{A}$ can be used to $(\epsilon c', \tau + c)$-find function collisions where c, c' are constants.*

Security of KSN in fact requires collision resistance of gSWIFFT function when the input is a vector of short length that is infinity norm is $\leqslant \frac{t}{2}$ and then the collision resistance can be reduced to finding a Short Integer Solution problem [3].

Definition 3.3. **Strong Unforgeability.** Let $\mathcal{A}$ has queried on a message M during the query phase of the security game and σ be the signature of the message. In the strong unforgeability game, $\mathcal{A}$ is considered to be successful even if it can return a signature σ^* of the message M where $\sigma^* \neq \sigma$.

SWIFFT hash function used in KSN-OTS is itself highly parallelizable and this leads to higher speed of K2SN-MSS compared to XMSS that is based on W-OTS$^+$ that uses l random walks in the function family $\mathcal{F}_n$ which are inherently sequential. Comparison of the two signature schemes is given in Section 6.2.

3.2 K2SN-MSS construction

K2SN-MSS extends the KSN-OTS to multi-message signature scheme and uses SWIFFT as the underlying hash function. However, SWIFFT compresses $2n$ bits to $n + n_\epsilon$ bits and so concatenation of two outputs (in a node of the Merkle tree) results in $2n + 2n_\epsilon$ bits. We introduce a function Merge that effectively slides one string over the other and XORs the overlapping part, to generate a $2n$-bit string. The overlapping part of the two strings is XOR-ed. For our chosen parameters $n = 512$ and $n_\epsilon = 64$.

$$\begin{aligned}
\text{Merge}(\mathbf{y}_1, \mathbf{y}_2) &= \text{hi}_{n-n_\epsilon}(\mathbf{y}_1) \| (\text{low}_{2n_\epsilon}(\mathbf{y}_1) \oplus \text{hi}_{2n_\epsilon}(\mathbf{y}_2)) \\
&\quad \| \text{low}_{n-n_\epsilon}(\mathbf{y}_2), \quad (7)
\end{aligned}$$

where (i) $\mathbf{y}_1$ and $\mathbf{y}_2$ are $(n + n_\epsilon)$-bit strings, (ii) $\text{hi}_l(\mathbf{y})$ is the most significant l bits of $\mathbf{y}$ and (iii) $\text{low}_l(\mathbf{y})$ is the least significant l bits of $\mathbf{y}$. The function Merge is used to calculate $y'_{j,i}$ Figure 2 instead of simple concatenation.

Lemma 3.4. *Let f be a random oracle which outputs $(n + n_\epsilon)$-bit long random strings. Let O_c and O_m be two oracles defined as follows: On a query, O_c outputs a $2n$-bit long string y which is computed as $y = \text{low}_n(x_1) \| \text{low}_n(x_2) : x_1, x_2 \xleftarrow{\$} f$. On the other hand, O_m outputs a $2n$-bit long string y' which is computed as $y' = \text{Merge}(x_1, x_2) : x_1, x_2 \xleftarrow{\$} f$. Then O_c and O_m are indistinguishable.*

Proof is in full version [25].

3.2.1 Algorithms. K2SN-MSS consists of three algorithms: $\texttt{K2SN.KeyGen}$, $\texttt{K2SN.Sign}$ and $\texttt{K2SN.Verify}$.

Key Generation: $\texttt{K2SN.KeyGen}$ uses two function families: F_n and SWIFFT as given in Algorithm 4. Here n is the length of the seed of the PRG that is used for secret key generation and is the security parameter of K2SN-MSS. During computation of each node in the tree a different hash key has been used. The public key length is $n + n_\epsilon + \texttt{hkseed} + \texttt{rpseed}$ bits.

Algorithm 4 Key Generation of K2SN-MSS ($\texttt{K2SN.KeyGen}(1^n)$)

(1) Choose randomly, (i) $\text{sk} \in \{0, 1\}^n$, (ii) $\texttt{hkseed} \in \{0, 1\}^n$ and (iii) $\texttt{rpseed} \in \{0, 1\}^n$.
(2) For $i = 0, 1, \ldots, T - 1$ (where $T = 2^h$) /* Construction of $\mathcal{L}_i$ trees */
   (a) For each $\mathcal{L}_i$ tree, compute $\text{sk}_i = \text{low}_n(f_{\text{sk}}(i + 1))$.
   (b) For $j = 0, 1, \ldots, t - 1$, ($t = 2^\ell$), compute secret key components of the i-th signature as, $\mathbf{x}_{i_j} = f_{\text{sk}_i}(j + 1)$.
   (c) For $j = 0, 1, \ldots, t - 1$, compute the public key component of i-th signature as, $\mathbf{y}_{i_j} = \text{SWIFFT}_{k_i}(\mathbf{x}_{i_j})$.
   (d) Compute $\mathbf{y}_{h,i}$, the root of the Merkle tree $\mathcal{L}_i$ from leaves $\mathbf{y}_{i_0}, \mathbf{y}_{i_1}, \ldots, \mathbf{y}_{i_{t-1}}$. All the remaining nodes of the $\mathcal{L}_i$ tree are computed with a new random pad and hash key.
(3) The roots of the $\mathcal{L}_i$ trees are the leaf nodes of the Merkle tree $\mathcal{MSS}$ at height h Denote the nodes by $\mathbf{y}_{h,0}, \mathbf{y}_{h,1}, \ldots, \mathbf{y}_{h,T-1}$, and construct the Merkle tree $\mathcal{MSS}$ using different hash keys and random pads. The root of the tree is denoted by $\mathbf{y}_{0,0}$.
(4) Return, (i) Secret key $\mathcal{SK} = \text{sk}$, and (ii) Public key $\mathcal{PK} = \{\mathbf{y}_{0,0}, \texttt{hkseed}, \texttt{rpseed}\}$.

Signature Generation. For the i-th message, mes, the signature $\sigma_{\text{mes}} = (i, \text{pk}, \mathcal{PK}_i, \text{Auth})$, where $i, 1 \leq i \leq 2^h$ is the index of the signature, pk is the sum of the components secret key associated with B_{mes}, $\mathcal{PK}_i$ is the set of component public keys of the i-th signature (i-th instance of KSN-OTS) and Auth contains the nodes of the $\mathcal{MSS}$ tree that are required for verification of $\mathcal{L}_i$ tree. The computations of pk, $\mathcal{PK}_i$ and Auth are given in Algorithm 5. The signature length consists of h-bits to represent i, $(\lfloor \log_2(\frac{t}{2}) \rfloor + 1) \times 2n$ bits to represent pk, $\mathcal{PK}_i$ consisting of t outputs of SWIFFT resulting in $t(n + n_\epsilon)$ bits, and $h(n + n_\epsilon)$ bits to represent Auth which consists of the h sibling nodes. In total signature size is $|\sigma_{\text{mes}}| = h + n \times (2(\lfloor \log_2(\frac{t}{2}) \rfloor + 1) + t + h) + n_\epsilon \times (t + h)$ bits.

Signature Verification. To verify $\sigma_{\text{mes}} = (i, \text{pk}, \mathcal{PK}_i, \text{Auth})$, first pk is verified against $\mathcal{PK}_i$ using the KSN-OTS verification algorithm (Algorithm 3), and then $\mathcal{PK}_i$ is authenticated against $\mathbf{y}_{0,0}$, the root of the $\mathcal{MSS}$ Merkle tree using Auth. The algorithm is given in Algorithm 6.

Algorithm 5 Signature Generation of K2SN-MSS ($\texttt{K2SN.Sign}(i, \text{mes}, \mathcal{SK})$)

(1) Compute the seed sk_i for the $\mathcal{L}_i$ tree as $\text{sk}_i = f_{\text{sk}}(i+1)$.
(2) Compute $B_{\text{mes}} = c_m(\text{mes})$ and let $B_{\text{mes}} = \{i_0, i_1, \ldots, i_{\frac{t}{2}-1}\}$ such that $i_j < i_{j+1}$ for all $0 \leq j \leq \frac{t}{2} - 2$.
(3) For $j = 0, 1, \ldots, t-1$, compute the component secret keys of the i-th KSN-OTS as $\mathbf{x}_{i_j} = f_{\text{sk}_i}(j+1)$.
(4) Compute the component public keys of the i-th KSN-OTS as $\mathcal{PK}_i = \{\mathbf{y}_{i_j} = \text{SWIFFT}_k(\mathbf{x}_{i_j}), j = 0, 1, \ldots, t-1\}$.
(5) $\text{pk} = \mathbf{x}_{i_0} \oplus_p \mathbf{x}_{i_1} \oplus_p \cdots \oplus_p \mathbf{x}_{i_{\frac{t}{2}-1}}$, where $i_j \in B_{\text{mes}}$.
(6) Let $\mathbf{y}_{h,i}$ be the root of the Merkle tree $\mathcal{L}_i$. The Auth contains the sibling nodes of the $\mathcal{MSS}$ Merkle tree from $\mathbf{y}_{h,i}$ to the root $\mathbf{y}_{0,0}$.
(7) Return $\sigma_{\text{mes}} = (i, \text{pk}, \mathcal{PK}_i, \text{Auth})$.

Algorithm 6 Signature Verification of K2SN-MSS ($\texttt{K2SN.Verify}(\text{mes}, \sigma_{\text{mes}})$)

(1) Let $\sigma_{\text{mes}} = (i, \text{pk}, \mathcal{PK}'_i, \text{Auth})$.
(2) Compute $B_{\text{mes}} = c_m(\text{mes})$ and and let $B_{\text{mes}} = \{j_0, j_1, \ldots, j_{\frac{t}{2}-1}\}$ such that $j_k < j_{k+1}$ for all $0 \leq k \leq \frac{t}{2} - 2$.
(3) If $\text{pk}[i] \leq \frac{t}{2}, \forall 0 \leq i < \hat{m}\hat{n}$, then Continue else Reject.
(4) Let $\mathcal{PK}'_i = \{\mathbf{y}'_0, \mathbf{y}'_1, \ldots, \mathbf{y}'_{t-1}\} \in \sigma_{\text{mes}}$.
(5) Compute $\text{pk}' = \mathbf{y}'_{j_0} + \mathbf{y}'_{j_1} + \cdots + \mathbf{y}'_{j_{\frac{t}{2}-1}}$, where the additions are in the ring $\mathcal{R}$.
(6) If $\text{pk}' = g\text{SWIFFT}_k(\text{pk})$, then Continue, else Reject.
(7) Compute the root of the Merkle tree $\mathcal{L}_i$ where the leaf nodes are $\mathbf{y}'_0, \mathbf{y}'_1, \ldots, \mathbf{y}'_{t-1}$ and with the hash keys and random pads used during signing. The hash keys and random pads can be easily regenerated from the published seed. Let the root of the Merkle tree $\mathcal{L}_i$ be $\mathbf{y}'_{h,i}$.
(8) Compute the root of the $\mathcal{MSS}$ Merkle tree from $\mathbf{y}'_{h,i}$ using Auth as

$$\mathbf{y}'_{l, \frac{i}{2}} = \text{SWIFFT}_{k_{l, \frac{i}{2}}}(\text{Merge}(\mathbf{y}'_{l+1,i}, \text{Auth}[h-l]) \oplus \mathbf{v}_{l, \frac{i}{2}}),$$
$$\text{where } i = 0 \pmod 2,$$

$$\mathbf{y}'_{l, \frac{i-1}{2}} = \text{SWIFFT}_{k_{l, \frac{i-1}{2}}}(\text{Merge}(\text{Auth}[h-l], \mathbf{y}'_{l+1,i}) \oplus \mathbf{v}_{l, \frac{i-1}{2}}),$$
$$\text{where } i = 1 \pmod 2,$$

where $\text{Auth}[h-l]$ denotes the $(h-l)$-th node of the Auth.
(9) If $\mathbf{y}'_{0,0} = \mathbf{y}_{0,0}$ (where $\mathbf{y}_{0,0}$ is the published root of the $\mathcal{MSS}$ Merkle tree as a part of $\mathcal{PK}$), then Accept, else Reject.

3.3 K2SN-MSS Security

THEOREM 3.5. *If $\mathcal{H}$ is a $2^h(2^\ell - 1)$ multi-function multi-target second-preimage resistant hash function then K2SN-MSS is $(\epsilon, 2^h, \tau)$ secure against Strong Existential Unforgeable under Chosen Message attack (SUF-CMA). This means that for any forger $\mathcal{A}$ with success probability ϵ against K2SN-MSS after 2^h query in time τ. The upper*

bound of ϵ is

$$\epsilon \leq \frac{InSec^{\text{MM-SPR}}(\mathcal{H}, 2^h, 2^h(1+2^\ell))}{\epsilon'}, \text{ where}$$

$$\epsilon' = (1 - 2^{-2n} - 2^{-m})(1 - 2^{-2n}) \cdot \max\left\{2^{-2h}2^{-2(h+\ell)}, (1-2^{-2h})2^{-2(h+\ell)}\right\}.$$

where $q_1, q \leq 2^h$, and $q_2 \leq 2^{h+\ell} - 1$ respectively.

PROOF. The sketch of the proof is given in Appendix B and we refer the full version [25] for the detailed proof.

$\square$

Estimating security level of K2SN-MSS. In [21] it was argued that (See Equation (16) of [21]), if a hash function $\mathcal{H}$ is considered as random oracle, then

$$InSec^{\text{MM-SPR}}(\mathcal{H}, q, q') = \frac{(q+1)q'}{2^{n'}},$$

where $\mathcal{H}_K : \{0,1\}^{2n} \mapsto \{0,1\}^{n'}$.

In K2SN-MSS, we instantiate $\mathcal{H}$ by SWIFFT with parameters $\hat{m} = 16$, $\hat{n} = 64$ and $p = 257$. Therefore, n', the size of the hash output is 576 bits.

For $h = 20$, $\ell = 9$, $m = 256$, $n = 512$ and $n' = 576$, we have $(1 - 2^{-2n} - 2^{-m})(1 - 2^{-2n}) \approx 1$, and so,

$$\epsilon \leq \frac{2^{40}(1+2^9)2^{-576}}{\max\{2^{-40}2^{-58}, (1-2^{-40})2^{-58}\}} \approx 2^{-469}.$$

The bit-security of the signature scheme is computed as $\log_2(\epsilon/\tau)$ [7], where τ is the time (measured in the number of hash evaluation) required for key generation, q signature generations, q verifications and the time required for $\mathcal{A}$ to forge a signature [42]. Key generation requires $2^{h+\ell+1} - 1$ computations of SWIFFT, signing 2^h signatures requires $2^h \times (h+1) \times 2^{\ell+1}$ hash computations and verification of 2^h signatures requires $2^h \times (h + 2^{\ell+1})$ hash computations [11]. Therefore, the lower bound on τ is

$$\tau \geq 2^{h+\ell+1} + 2^h \times (h+1) \times 2^{\ell+1} + 2^h \times (h + 2^{\ell+1}) \approx 2^{35}$$

and the bit-security of K2SN-MSS for the given parameters will be 504.

4 SIMD PARALLELIZATION OF NTT

Our implementation of K2SN-MSS uses SIMD parallelization of SWIFFT. In the following we show how ring multiplication, which is the most costly operation in SWIFFT, can be parallelized using SIMD instructions of avx2. We use the level of parallelization as a parameter that can be chosen based on the setting.

Notations. We will denote a vector by bold lowercase alphabet (e.g. $\mathbf{z}, \boldsymbol{\beta}$). By a single subscript, we denote a vector of dimension 64 (e.g. $\mathbf{z}_i$ is a vector of length 64). If we use two subscript separated by a ",", then it denotes a vector of length 2^δ. We denote the k-th component of a vector $\mathbf{z}_i$ (or $\mathbf{z}_{i,j}$) as $\mathbf{z}_i[k]$ (or $\mathbf{z}_{i,j}[k]$). We use $\oplus_p$ and $\odot_p$ as defined in Section 2. The main computation of K2SN-MSS is the computation of SWIFFT. We describe our SIMD parallelization of this computation for the following parameters:

$$\hat{n} = 64, \hat{m} = 16, p = 257 \text{ and } \omega = 42 \pmod{257},$$

These are the parameters given in [2, 30]. For these parameters, the output of SWIFFT and gSWIFFT are elements of $\mathbb{Z}_{257}^{64}$. The input to

SWIFFT is from $\mathbb{Z}_2^{1024}$ and to gSWIFFT is from $\mathbb{Z}_{257}^{1024}$. An input vector $\mathbf{z}$ is mapped into elements of $\mathcal{R}$ by partitioning the vector $\mathbf{z}$ into 16 sub-vectors of dimension 64 each, such that $\mathbf{z} = (\mathbf{z}_0, \mathbf{z}_1, \ldots, \mathbf{z}_{15})$, where each $\mathbf{z}_i$ can be represented as a ring element. Further assume that each $\mathbf{z}_i = (z_{i,0}, z_{i,1}, \ldots, z_{i,63})$, with $z_{i,j} \in \mathbb{Z}_2$ (or $z_{i,j} \in \mathbb{Z}_{257}$). Let $\mathbf{a}_0, \mathbf{a}_1, \ldots, \mathbf{a}_{\hat{m}-1}$ be the multipliers of a SWIFFT function. Then the hash of $\mathbf{z}$ is computed by equation (4) as $\sum_{i=0}^{15} \mathbf{a}_i \cdot \mathbf{z}_i$. This computation in practice is by equation (5) as $\sum_{i=0}^{15} \text{NTT}(\mathbf{a}_i) \odot_p \text{NTT}(\mathbf{z}_i)$ and so the main computation here is NTT.

Let $\mathbf{y}_i$ be the NTT of $\mathbf{z}_i$ given by $\mathbf{y}_i = \text{NTT}(\mathbf{z}_i) = (y_{i,0}, y_{i,1}, \ldots, y_{i,63})$. Our objective is to parallelize the computation of $\mathbf{y}_i = \text{NTT}(\mathbf{z}_i)$ for $0 \leqslant i \leqslant 15$. We introduce a parameter δ that determines the level of parallelization: that is a δ-parallelized computation of NTT, performs operation on 2^δ components of the vector $\mathbf{y}_i$. We divide $\mathbf{y}_i$ into $\frac{64}{2^\delta}$ vectors of dimension 2^δ such that $\mathbf{y}_i = \left(\mathbf{y}_{i,0}, \mathbf{y}_{i,1}, \ldots, \mathbf{y}_{i,\frac{64}{2^\delta}-1} \right)$, where $\mathbf{y}_{i,\ell} = \left(y_{i,2^\delta \cdot \ell + 0}, y_{i,2^\delta \cdot \ell + 1}, \ldots, y_{i,2^\delta \cdot \ell + 2^\delta - 1} \right)$ for $0 \leqslant \ell \leqslant \frac{64}{2^\delta} - 1$. Therefore, each $y_{i,j}$ is j_0^{th} component of the vector $\mathbf{y}_{i,j_1}$ for some $0 \leqslant j_0 \leqslant 2^\delta - 1$ and $0 \leqslant j_1 \leqslant \frac{64}{2^\delta} - 1$. We rewrite the j of equation (8) as $j = j_0 + 2^\delta j_1$, where $0 \leqslant j_0 \leqslant 2^\delta - 1$ and $0 \leqslant j_1 \leqslant \frac{64}{2^\delta} - 1$. As a consequence, k of the equation (8) has to broken down into k_0 and k_1 as $k = k_0 + \frac{64}{2^\delta} k_1$, where $0 \leqslant k_0 \leqslant \frac{64}{2^\delta} - 1$ and $0 \leqslant k_1 \leqslant 2^\delta - 1$. Now we compute each $y_{i,j}$ as,

$$y_{i,j} = \sum_{k=0}^{63} (z_{i,k} \omega^k) \omega^{2jk}, \forall j, 0 \leqslant j \leqslant 63. \tag{8}$$

The equation (8) can be restated as:

$$y_{i,j_0+2^\delta j_1} = \sum_{k_0=0}^{\frac{64}{2^\delta}-1} \left(\omega^{2^{\delta+1}} \right)^{j_1 k_0} \cdot y'_{j_0,k_0}, \tag{9}$$

where

$$y'_{j_0,k_0} = \alpha_{j_0,k_0} \cdot \beta_{j_0,k_0}, \tag{10}$$

$$\alpha_{j_0,k_0} = \omega^{(2j_0+1)k_0}, \tag{11}$$

$$\beta_{j_0,k_0} = \sum_{k_1=0}^{2^\delta-1} \left(\omega^{\frac{64}{2^\delta} k_1 (1+2j_0)} z_{i,k_0+\frac{64}{2^\delta} k_1} \right). \tag{12}$$

For δ-parallelization, let $\mathbf{y}_{i,j_1} = (y_{i,2^\delta \cdot j_1 + 0}, y_{i,2^\delta \cdot j_1 + 1}, \ldots, y_{i,2^\delta \cdot j_1 + 2^\delta - 1})$. We compute the components of the vectors $\mathbf{y}_{i,j_1}$ in parallel for $0 \leqslant j_1 \leqslant \frac{64}{2^\delta} - 1$. From equation (9), we know $y_{i,2^\delta \cdot j_1 + j_0}$ (for $0 \leqslant j_0 \leqslant 2^\delta - 1$) can be expressed as $\sum_{k_0=0}^{\frac{64}{2^\delta}-1} \left(\omega^{2^{\delta+1}} \right)^{j_1 k_0} \cdot y'_{j_0,k_0}$. Therefore, the vector $\mathbf{y}_{i,j_1}$ can be written as

$$\mathbf{y}_{i,j_1} = \left(\sum_{k_0=0}^{\frac{64}{2^\delta}-1} \left(\omega^{2^{\delta+1}} \right)^{j_1 k_0} \cdot y'_{0,k_0}, \ldots, \sum_{k_0=0}^{\frac{64}{2^\delta}-1} \left(\omega^{2^{\delta+1}} \right)^{j_1 k_0} \cdot y'_{2^\delta-1,k_0} \right). \tag{13}$$

Thus the vector $\mathbf{y}_{i,j_1}$ can be computed from the vector $\mathbf{y}'_{j_0,k_0} = (y'_{0,k_0}, y'_{1,k_0}, \ldots, y'_{2^\delta-1,k_0})$ by multiplications by scalars $(\omega^{2^{\delta+1}})^{j_1 k_0}$, $0 \leqslant k_0 \leqslant \frac{64}{2^\delta} - 1$.

On the other hand, each component y'_{j_0,k_0} of the vector $\mathbf{y}'_{j_0,k_0}$ depends on α_{j_0,k_0} and β_{j_0,k_0} through the equation (10). Thus, we can express the vector $\mathbf{y}'_{j_0,k_0}$ as $\mathbf{y}'_{j_0,k_0} = (\alpha_{0,k_0} \cdot \beta_{0,k_0}, \alpha_{1,k_0} \cdot \beta_{1,k_0}, \ldots, \alpha_{2^\delta-1,k_0} \cdot \beta_{2^\delta-1,k_0})$, which can be viewed as the component-wise

multiplication of the vectors $\boldsymbol{\alpha}_{j_0,k_0} = (\alpha_{0,k_0}, \alpha_{1,k_0}, \ldots, \alpha_{2^\delta-1,k_0})$, and $\boldsymbol{\beta}_{j_0,k_0} = (\beta_{0,k_0}, \beta_{1,k_0}, \ldots, \beta_{2^\delta-1,k_0})$.

The vectors $\boldsymbol{\alpha}_{j_0,k_0}$ can be precomputed as they do not depend on the input string. If $\mathbf{z} \in \mathbb{Z}_2^{1024}$, the vectors $\boldsymbol{\beta}_{j_0,k_0}$ can be precomputed or computed in real-time during the execution and this depends on the level of parallelization. To compute vectors $\mathbf{y}_i$ in parallel, we first compute vectors $\mathbf{y}'_{j_0,k_0}$ from $\boldsymbol{\alpha}_{j_0,k_0}$, and $\boldsymbol{\beta}_{j_0,k_0}$, and then using equation (13), compute the vectors $\mathbf{y}_{i,j_1}$ from the vectors $\mathbf{y}'_{j_0,k_0}$. Details are in Section 4.1.

4.1 Parallelization of NTT when $\mathbf{z}_i \in \mathbb{Z}_2^{64}$ ($b\text{NTT}2^\delta(\cdot)$)

For $\mathbf{z}_i \in \mathbb{Z}_2^{64}$, the computation is called $b\text{NTT} - 2^\delta(\cdot)$ and given below.

(1) **Precomputation of vectors $\boldsymbol{\alpha}_{j_0,k_0}$:** Define vectors $\boldsymbol{\alpha}_{j_0,k_0} = (\alpha_{0,k_0}, \alpha_{1,k_0}, \ldots, \alpha_{2^\delta-1,k_0}) \in \mathbb{Z}_{257}^{2^\delta}$, where α_{j_0,k_0} is defined by equation (11) for $0 \leqslant k_0 \leqslant \frac{64}{2^\delta} - 1$. The vectors $\boldsymbol{\alpha}_{j_0,k_0}$ can have $\frac{64}{2^\delta}$ possible values depending on the value of k_0. Let Λ_1 be the precomputation table where rows of the table are the vectors $\boldsymbol{\alpha}_{j_0,k_0}$. There are k_0 rows of the table Λ_1 which are indexed by k_0.

(2) **Precomputation of vectors $\boldsymbol{\beta}_{j_0,k_0}$:** Define vectors $\boldsymbol{\beta}_{j_0,k_0} = (\beta_{0,k_0}, \beta_{1,k_0}, \ldots, \beta_{2^\delta-1,k_0}) \in \mathbb{Z}_{257}^{2^\delta}$, where β_{j_0,k_0} is defined by equation (12). For a given k_0, $\boldsymbol{\beta}_{j_0,k_0}$ can have 2^{2^δ} possible values depending on the combination of the bits

$$\mathbf{z}_{i,k_0} = \left(z_{i,k_0+\frac{64}{2^\delta} \cdot 0}, z_{i,k_0+\frac{64}{2^\delta} \cdot 1}, \ldots, z_{i,k_0+\frac{64}{2^\delta} \cdot (2^\delta-1)} \right), \tag{14}$$

that can be easily precomputed and stored. Let the precomputation table Λ_2 with these entries be indexed from 0 to $2^{2^\delta} - 1$. We can choose an entry from the table by the decimal value of $\mathbf{z}_{i,k_0}$ computed from the expression,

$$\text{dec}_{2^\delta}(\mathbf{z}_{i,k_0}) = z_{i,k_0+\frac{64}{2^\delta} \cdot 0} 2^0 + \cdots + z_{i,k_0+\frac{64}{2^\delta} \cdot (2^\delta-1)} 2^{2^\delta-1}.$$

(3) **Computation of vectors $\mathbf{y}'_{j_0,k_0}$:** For $0 \leqslant j_0 \leqslant 2^\delta - 1$, the j_0^{th} component of $\mathbf{y}_{j_0,k_0}$ only depends on the j_0^{th}-components of the elements $\Lambda_1(k_0)$s and $\Lambda_2(\text{dec}_{2^\delta}(\mathbf{z}_{i,k_0}))$s of $0 \leqslant k_0 \leqslant \frac{64}{2^\delta} - 1$. Following the observation, we can compute, for $0 \leqslant k_0 \leqslant \frac{64}{2^\delta} - 1$,

$$\mathbf{y}'_{j_0,k_0} = \Lambda_1(k_0) \odot_{257} \Lambda_2(\text{dec}(\mathbf{z}_{i,k_0})). \tag{15}$$

(4) **Computation of the output vectors $\mathbf{y}_{i,j_1}$:** Using the vectors $\mathbf{y}'_{j_0,k_0}$, $0 \leqslant k_0 \leqslant \frac{64}{2^\delta} - 1$, the output vectors $\mathbf{y}_{i,j_1}$ can be computed as:

$$\mathbf{y}_{i,j_1} = \sum_{k_0=0}^{\frac{64}{2^\delta}-1} (\omega^{16})^{j_1 k_0} \mathbf{y}'_{j_0,k_0}, \forall 0 \leqslant j_1 \leqslant \frac{64}{2^\delta} - 1. \tag{16}$$

In order to compute the vectors $\mathbf{y}_{i,j_1}$ of the equation (16), 2^δ vector multiplications are required and there are $\frac{64}{2^\delta}$ such vectors. Therefore, the computation is 64 vector multiplications.

Memory requirement of the precomputation: Required memory for the table Λ_1 of vectors $\boldsymbol{\alpha}_{j_0,k_0}$ is of $O\left(\frac{64}{2^\delta}\right)$. The size of

the table Λ_2 of vectors $\boldsymbol{\beta}_{j_0, k_0}$ is of order $O\left(2^{2^{\delta}}\right)$. The memory requirement is thus dominated by the size of the Λ_2 and it increases exponentially as δ increases.

4.2 Parallelization of NTT when $\mathbf{z}_i \in \mathbb{Z}_{257}^{64}$ (gNTT$2^{\delta}(\cdot)$)

The input to the function gNTT is a vector in $\mathbb{Z}_{257}^{64}$. The computation of $\mathbf{y_i} = g\text{NTT}(\mathbf{z}_i)$ is the same as the computation of the function bNTT except for Step 2, for which we provide details below.

The components of $\mathbf{z}_{i, k_0}$ in equation (14) are now elements from $\mathbb{Z}_{257}$ and so vector $\mathbf{z}_{i, k_0} = \left(z_{i, k_0 + \frac{64}{2^{\delta}} \cdot 0}, z_{i, k_0 + \frac{64}{2^{\delta}} \cdot 1}, \ldots, z_{i, k_0 + \frac{64}{2^{\delta}} \cdot (2^{\delta} - 1)}\right)$ has $257^{2^{\delta}}$ possible values. If we processed the same as the computation of bNTT, the size of the table Λ_2 will be $2^{2^{3\delta}}$ which is not acceptable. Therefore, we compute the required $\boldsymbol{\beta}_{j_0, k_0}$ vectors during the computation of gNTT$(\cdot)$ as given in Algorithm 7.

Define function $\text{Select2}^{\delta}(\mathbf{z}_{i, k_0}, j) = (\mathbf{z}_{i, k_0}[j], \mathbf{z}_{i, k_0}[j], \ldots, \mathbf{z}_{i, k_0}[j])$. To compute vectors $\boldsymbol{\beta}_{j_0, k_0}$ as in Algorithm 7, we also require the following precomputed vectors

$$\boldsymbol{\omega}_{k_1, j_0} = \left(\omega^{\frac{64}{2^{\delta}} k_1(1 + 2 \cdot 0)}, \omega^{\frac{64}{2^{\delta}} k_1(1 + 2 \cdot 1)}, \ldots, \omega^{\frac{64}{2^{\delta}} k_1(1 + 2 \cdot (2^{\delta} - 1))}\right)$$
$$\text{for } 0 \leq k_1 \leq 2^{\delta} - 1.$$

Algorithm 7 Parallelized computation of $\boldsymbol{\beta}_{j_0, k_0}$

Input: $(\mathbf{z}_{i,0}, \ldots, \mathbf{z}_{i, \frac{64}{2^{\delta}} - 1})$ and $(\boldsymbol{\omega}_{0, j_0}, \ldots, \boldsymbol{\omega}_{2^{\delta} - 1, j_0})$;

Output: $\boldsymbol{\beta}_{j_0, 0}, \boldsymbol{\beta}_{j_0, 1}, \ldots, \boldsymbol{\beta}_{j_0, \frac{64}{2^{\delta}} - 1}$;

$\boldsymbol{\beta}_{j_0, k_0} = \boldsymbol{\beta}_{j_0, k_0} \oplus_{257} \left(\boldsymbol{\omega}_{k_1, j_0} \odot_{257} \text{Select2}^{\delta}(\mathbf{z}_{i, k_0}, k_1)\right), \forall 0 \leq k_0 \leq \frac{64}{2^{\delta}} - 1, \forall 0 \leq k_1 \leq 2^{\delta} - 1$

4.3 Optimization of the Step-4 of bNTT2^{δ} and gNTT2^{δ} for $\delta = 3$

To compute vectors $\mathbf{y}_{i, j_1}$ of the equation (16), 64 vector multiplications are required. However, the particular choice of $\omega \equiv 42 \pmod{257}$, we have $\omega^{16} \equiv 2^2 \pmod{257}$ and $(\omega^{16})^4 \equiv -1 \pmod{257}$. This will remove the need for multiplication and simplify the computations of $\mathbf{y}_{i, j_1}$s of equation (16) using left shifts by 2 and addition operations only. Table 5 (See Appendix C) includes the computations of the $\mathbf{y}_{i, j_1}$ and the expressions that involve only component-wise addition or subtraction. In Table 6 (See Appendix C), we list all the sub-expressions which we have to compute to obtain the final expressions of Table 5. The expressions of the Table 6 include component-wise addition or subtraction, left shifts by 2, 4 and 6 only.

4.4 Optimization of the Step-4 of bNTT2^{δ} and gNTT2^{δ} for $\delta = 4$

For $\omega \equiv 42 \pmod{257}$, we have $\omega^{32} \equiv 2^4 \pmod{257}$ and $(\omega^{32})^2 \equiv -1 \pmod{257}$. We can simplify the computations of $\mathbf{y}_{i, j_1}$s of equation (16) using left shifts by 4 and addition/subtraction operations

only, without any multiplication. Table 7 includes all the computations of the $\mathbf{y}_{i, j_1}$ and its all the expressions involve only component-wise addition or subtraction. On the other hand, Table 8 (See Appendix C) includes all the sub-expressions which are required to compute the final expressions of Table 7. Observe that the expressions of the Table 8 (See Appendix C) only need component-wise addition and subtraction and left shifts by 4.

4.5 Parallelizing SWIFFT with parameter δ

Let $(\mathbf{z}_0, \ldots, \mathbf{z}_{15})$ and $\mathbf{x}_S = (\mathbf{x}_{S, 0}, \ldots, \mathbf{x}_{S, \frac{64}{2^{\delta}} - 1}), \mathbf{x}_{S, i} \in \mathbf{Z}_{257}^{2^{\delta}}, \forall i$, be the input and output of SWIFFT, respectively, and $(\mathbf{a}_0, \mathbf{a}_1, \ldots, \mathbf{a}_{15})$ denote the vector of multipliers. The pre-computation for SWIFFT and gSWIFFT, both, will be for gNTT$2^{\delta}(\mathbf{a}_i) = (\mathbf{A}_{i, 0}, \mathbf{A}_{i, 1}, \ldots, \mathbf{A}_{i, \frac{64}{2^{\delta}}})$. The computation of $\mathbf{x}_S$ is given in the Algorithm 8.

Algorithm 8 SWIFFT-2^{δ}

Input: $(\mathbf{z}_0, \ldots, \mathbf{z}_{\hat{m}-1}), \mathbf{z}_i \in \mathbb{Z}_2^{64}; (\mathbf{A}_{i, 0}, \ldots, \mathbf{A}_{i, \frac{64}{2^{\delta}}}), 0 \leq i \leq \hat{m} - 1;$

Output: $\mathbf{x}_S = (\mathbf{x}_{S, 0}, \ldots, \mathbf{x}_{S, \frac{64}{2^{\delta}} - 1}) \in \mathbb{Z}_{257}^{64};$

---

(1) Compute $\mathbf{y}_i = b\text{NTT}2^{\delta}(\mathbf{z}_i) = (\mathbf{y}_{i, 0}, \ldots, \mathbf{y}_{i, \frac{64}{2^{\delta}} - 1}), \forall 0 \leq i \leq \hat{m} - 1;$

(2) Compute $\mathbf{t}_{i, j} = \mathbf{y}_{i, j} \odot_{257} \mathbf{A}_{i, 1}, \forall 0 \leq i \leq \hat{m} - 1, \forall 0 \leq j \leq \frac{64}{2^{\delta}}.$

(3) $\mathbf{x}_{S, i} = \mathbf{x}_{S, i} \oplus_{257} \mathbf{t}_{j, i}, \forall 0 \leq i \leq \frac{64}{2^{\delta}}, \forall 0 \leq j \leq \hat{m} - 1.$

In gSWIFFT-2^{δ}, the input vector $\mathbf{z}_i$s are in $\mathbb{Z}_{257}^{64}$, and the computation is given by Algorithm 9 (Appendix A), where the inputs are processed by the function gNTT2^{δ} while the rest of the computation is the same as in Algorithm 8.

Note that NTT computations and SWIFFT function evaluation only use vector operations $\oplus_{257}$ and $\odot_{257}$ that are performed on the vectors component-wise. The output vector of these vector operations can be directly used for the next vector operation. Thus the above 2^{δ}-way parallelization of NTT implementation of SWIFFT function is a suitable candidate for SIMD-based parallelization. The choice of δ depends on the available intrinsic operations that are required for integer SIMD operations.

We provide two new implementations of SWIFFT for $\delta = 3$ and $\delta = 4$ using avx2 intrinsic [22], supporting 256-bit registers for integer operations. We refer to the two implementations as SWIFFT-8-avx2 and SWIFFT-16-avx2, achieving 8-way and 16-way parallelizations, respectively. SWIFFT-16-avx2 has the highest parallelization level to the date which gives the fastest implementation. We obtained better implementation results for K2SN-MSS using SWIFFT-16-avx2 than SWIFFT-8-avx2 and so in the following we give implementation details of SWIFFT-16-avx2 only. We omit the implementation details of SWIFFT-8-avx2 here due to page limits. All implemented codes are publicly available at GitHub [24].

The existing benchmark implementation of SWIFFT in [34] corresponds to [2, 30]. The implementation uses $\delta = 3$ and parallelizes the code using sse2 [22] instructions of Intel processors. We refer to this implementation as SWIFFT-8-sse2.

5 K2SN-MSS SOFTWARE IMPLEMENTATION

In this section, we provide the details of the implementation of the software K2SN-MSS. Table 2 lists the values of the parameters used for the software K2SN-MSS and also the functions are used.

Parameters		
Symbol	Value	Meaning
n	512	Security parameter
m	256	Message Space
$\hat{n}$	64	Dimension of the ring $\mathcal{R}$ of SWIFFT
$\hat{m}$	16	Number of multipliers of SWIFFT
p	257	Prime of SWIFFT
t	262	optimized CFF for 256-bit messages
Functions		
Name	Description	
Hash Function: SWIFFT	$\{0,1\}^{\hat{n}\hat{m}} \mapsto \{0,1\}^{\hat{n}\lceil \log_2(p) \rceil}$	
Hash Function: gSWIFFT	$\{0,1\}^{\hat{n}\hat{m}\lceil \log_2(p) \rceil} \mapsto \{0,1\}^{\hat{n}\lceil \log_2(p) \rceil}$	
PRF ChaCha20	$\{0,1\}^{\hat{m}\hat{n}} \mapsto \{0,1\}^{\hat{m}\hat{n}}$	
1CFF: c_m	$\{0,1\}^m \mapsto \{i_0, i_1, \ldots, i_{\frac{t}{2}-1}\}$,	
	where each $0 \leqslant i_j < t$ and $i_{j_1} \neq i_{j_2}$	
	for any $0 \leqslant j_1 < j_2 < \frac{t}{2}$	
Sizes		
Name	Lengths	
Signature Size	21331 Bytes	
Secret Key Size	40 Bytes	
Public Key Size	152 Bytes	

Table 2: K2SN-MSS parameters, functions and resulting sizes

The performance of K2SN-MSS is dominated by the computation of SWIFFT. We describe our SIMD parallelization of this computation for the following parameters which are given in [2, 30]:

$$\hat{n} = 64, \hat{m} = 16, p = 257 \text{ and } \omega = 42 \ (\text{mod } 257),$$

Our reasoning for these choices are below. Lyubashevsky et al. reduced the SWIFFT function to the subset sum problem and used k-list attack to compute the preimage of a SWIFFT function [30] in practice.

This choice of parameters also allows us to compare our implementation of SWIFFT against existing software of SWIFFT [34], both for correctness and efficiency. Note that in $\mathbb{Z}_p$, 42 is a $2\hat{n} = 128$-th root of unity. Therefore, for NTT, we use $\omega = 42 \ (\text{mod } 257)$ and show that with this choice of ω, multiplications reduce to bit-wise left-shift operations and results efficient implementation.

For $\hat{m} = 16$ and $\hat{n} = 64$, the input and output of SWIFFT are, 1024 and 576 bits, respectively. Therefore, SWIFFT is a hash function which compresses $2n$ bits to $n + n_\epsilon$ with $n = 512$ and $n_\epsilon = 64$.

Our implemented software is for 256-bit messages. Relaxing the condition $\frac{t}{2} < \frac{p}{2}$ by 3, we choose $\frac{t}{2} = 131$, such that t becomes the smallest even positive integer 262 where $\log_2 \binom{262}{131} \geqslant 256$. Therefore we use $t = 262$ for our software. We start with by providing the implementation details using intel avx2 intrinsic.

On the other hand, choice of these parameters for SWIFFT hash function leads to 512-bit classical (256-bit quantum) security of K2SN-MSS in multi-function-multi-target model. Therefore, we can only compare our implemented signature scheme against hash-based signature schemes which provide same level of security in

the same security model. The most optimized implementation of XMSS at the same security level is given in [13]. It should be noted that there is no existing software of XMSS^{MT} and SPHINCS at the same security level.

5.1 SWIFFT-16-avx2

In the implementation of 16-way parallelization, each 256-bit register __m256i is partitioned into 16 blocks of 16-bit each. The input strings are over $\mathbb{Z}_{257}$, where each element is 9-bit. The resulting vector of 16 $\mathbb{Z}_{257}$-elements is stored in a register. This is called *packing*. After completing the computation, the register content is moved to an array of integers. This is called *unpacking*.

We will denote packed integers as vectors of $(\text{int}_0, \text{int}_1, \ldots, \text{int}_{15})$, each int_i up to 16-bit. Here, unless otherwise stated, vectors are of dimension 16. Initially, each of the 16-bit blocks contains an element of $\mathbb{Z}_{257}$ which is represented by 9-bit. The result of vector addition and multiplication must be reduced modulo $p = 257$. The Details of the avx2 implementations of modular reduction, and the three main vector operations addition, subtraction and multiplication are given in Appendix D and the code is publicly available at [24].

5.2 Implementation of Cover-Free-Family

KSN uses a 1-CFF which is obtained by taking all subsets of size $t/2$ from a set of size t and so one can use a $\log_2 \binom{t}{t/2}$ bit message space. Each subset corresponds to a particular message. For efficient encoding of messages however, we use the algorithm proposed in [6, 10, 42].

5.3 Implementation of Pseudo-Random Function

For proof of theorem we assume all the randomness are uniform. Then in implementation we use a pseudorandom generator. We use ChaCha20 [4] as the pseudo-random function family F_n. ChaCha20 is an state-of-the-art stream-cipher which we use to generate the seed of each KSN-OTS from secret-key of K2SN-MSS, and all the component secret keys of the OTS instances. Details of the hash keys and random pads generation are given in full version [25]. We use the avx2-based implementation of ChaCha20 from supercop [26][1].

6 EXPERIMENTS

We used the following platform for our implementations and experiments:

Skylake: Intel®Core™i7-6700 4-core CPU @ 3.40GHz running. The timing experiments are performed on a single core. The OS is 64-bit Ubuntu-18.04 LTS and C codes were compiled by GCC version 7.3.0. The code of the software is available at [24]. During the experiments, the turbo boost and hyper-threading were turned off. The cache warm-up was done by 25000 iterations and the measurements are taken as the median over 100000 iterations.

---

[1] The software is implemented by D. J. Bernstein and R. Dolbeau and available in the directory "supercop-20171218/crypto_stream/chacha20/dolbeau/amd64-avx2"

6.1 Performance Comparison of SWIFFT Implementations

We have implemented all the parallelized version of the SWIFFT function using 16-bit avx2 intrinsic instructions. Each element of $\mathbb{Z}_{257}$ is represented by 9-bit. The SWIFFT evaluations were computed over 1024-bit data blocks. The generalized SWIFFT function uses 9×1024 bit data blocks as input. The output in both cases is a $9 \times 64 = 576$ bits string. For binary versions of SWIFFT-8-avx2 and SWIFFT-16-avx2, we require approximately 8KB and 2MB memory for the precomputation tables, respectively. The Time Stamp Counter (TSC) was read from the CPU to RAX and RDX registers using RDTSC instruction. All the experimental results are listed in Table 3. The results show that SWIFFT-8-avx2 and SWIFFT-16-avx2 are approximately 8% and 25% faster than the previous implementation in [2, 30]. Based on these results, we use SWIFFT-16-avx2 for K2SN-MSS.

Function Name	Intrinsic Flag	Binary Version		Generalized Version	
		Total clk	clk/byte	Total Clk	clk/byte
SWIFFT-8-sse2 [34]	mavx2	1150	8.98	-	-
SWIFFT-8-avx2 [this paper]	mavx2	1047	8.17	11435	9.93
SWIFFT-16-avx2 [this paper]	mavx2	866	6.77	9535	8.27

Table 3: Required clock cycles (clk) for various SWIFFT implementation.

6.2 Performance Comparison of KSN and W-OTS$^+$

KSN-OTS and W-OTS$^+$ are the OTS that are used in K2SN-MSS and XMSS respectively. We compare both of them at $n = 512$. For the W-OTS$^+$, we used the code for XMSS [19], available at [17, 18]. The Time Stamp Counter reading did not work with the XMSS code of [17, 18]. We instead used clock() function of "time.h" header file. Our measurement is the average over 1,000,000 iterations. The results of performance comparison of KSN-OTS and W-OTS$^+$ are in Table 4 and it shows that, the key generation of KSN is approximately 22 times faster than that of W-OTS$^+$, while signing and verification are approximately 23 and 167 times faster than those of W-OTS$^+$. This performance is due to the simple signing operation (generation of component secret keys and component-wise vector addition) and efficient implementation of gSWIFFT that is used in the verification operation of KSN-OTS.

	KSN/SWIFFT-16-avx2	W-OTS$^+$/SHA512/$w = 16$
Key Generation (μs)	164	3575
Signing (μs)	83	1872
Verification (μs)	10	1674
Secret Key/ (Bytes)	40	64
Public Key (Bytes)	4608	4352
Signature Sizes (Bytes)	1024	4288

Table 4: Performance Comparison of KSN and W-OTS$^+$ in μs

Remark: We compare our implemented K2SN-MSS software against the results of [13], but the code of [13] is not public.[2] [13] reports the fastest result for XMSS.

---

[2]We tried to communicate the authors, but did not receive any reply yet.

6.3 Performance Comparison between K2SN-MSS and XMSS

The timing measurements are done the same as in Section 6.2. We compare our implemented software against the XMSS software results available at [13]. We used the SWIFFT-16-avx2 implementation of K2SN-MSS. We compute the authentication path using the TREEHASH algorithm of [8]. The results of the performance comparisons are in Table 1 and it shows that the key generation, signing and verification of K2SN-MSS are 2.76 times, 2.89 times and 2.65 times faster than the corresponding operations in XMSS [13], respectively.

Remarks:

(1) The signature size of K2SN-MSS is comparable with the signature sizes of the XMSS^{MT} [20] and the SPHINCS [5], but we can not compare K2SN against XMSS^{MT} and SPHINCS because the use of the Merkle tree is different in the two cases. In K2SN-MSS and XMSS (single tree version of XMSS^{MT}) use only one layer of Merkle tree where XMSS^{MT} and SPHINCS uses multiple layers of Merkle tree.

(2) Even so, if we want to compare K2SN-MSS against the XMSS^{MT} and SPHINCS, we need optimized software of them for $n = 512$. But in the literature, we could not find a single implementation of XMSS^{MT} and SPHINCS which provides 512-bit classical (256-bit quantum) security. Thus we are unable to compare them.

(3) Our comparison is against the multi-buffer entries of the third row of TABLE IV of [13]. It is due to the following reasons:

  (a) this is the only reported results for SHA512, where SWIFFT output is 576 bits. This comparison is fair because of the comparative sizes of the used hash functions.

  (b) The experimental platform of [13] was
  Skylake: Intel®Core™i7-6700 4-core CPU @ 4.0GHz which is a faster machine than the machine used in our experiments. Therefore, we believe that the comparisons are made in Table 1 are valid and fair.

7 CONCLUDING REMARKS

Hash-based signatures are a strong alternative for post-quantum signatures. We extended KSN-OTS for signing multiple messages which is secure under multi-function multi-target attacks and gave an efficient implementation using parallelization at the instruction level, using a widely accessible technology of Intel. Our implementation also improves the state-of-the-art implementation of SWIFFT and provides parallelization of NTT computation, both of independent interest. Although our results are for concrete levels of parallelization, but by providing implementation details we provide a template for other parameters, for binary and non-binary input vectors, and with or without precomputation. Our implementation shows that K2SN-MSS is significantly faster than XMSS, which is recently proposed as a candidate for standardization.

In our implementation we used SWIFFT both for KSN-OTS, and the Merkle hash tree. Although SWIFFT is essential in KSN-OTS, the Merkle tree is used to authenticate the public key $\mathcal{PK}_i$, against the root of the $\mathcal{MSS}$ tree and we are not restricted to SWIFFT for the hash function. Thus, for this tree one can use traditional hash

functions such as SHA512. Our choice of SWIFFT hash function family for the construction of $\mathcal{L}$ trees and MSS has the following advantages. Firstly, it reduces the complexity of the code: using two different hash function families will increase the code complexity and size. Secondly, SWIFFT has (asymptotic) provable security and this improves confidence in the security of the design. Finally, SHA512 hash function family follows Merkle-Damgård construction [41] which is inherently a sequential construction and so not easily amenable to parallelization. SWIFFT function family, however, is highly parallelizable, and can benefit from processor architecture with longer registers to achieve higher speed.

8 ACKNOWLEDGE

REFERENCES

[1] 2018. The Internet Engineering Task Force. https://tools.ietf.org/html/rfc8391.
[2] Yuriy Arbitman, Gil Dogon, Vadim Lyubashevsky, Chris Peikert, Daniele Micciancio, and Alon Rosen. 2008. SWIFFTX: A Proposal for the SHA-3 Standard. https://www.eecs.harvard.edu/~alon/PAPERS/lattices/swifftx.pdfl.
[3] Rouzbeh Behnia, Muslum O. Ozmen, Attila A. Yavuz, and Mike Rosulek. 2018. TACHYON: Fast Signatures from Compact Knapsack. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security (CCS '18)*. ACM, 1855–1867. https://doi.org/10.1145/3243734.3243819
[4] Daniel J. Bernstein. 2008. ChaCha, a variant of Salsa20.
[5] Daniel J. Bernstein, Daira Hopwood, Andreas Hülsing, Tanja Lange, Ruben Niederhagen, Louiza Papachristodoulou, Michael Schneider, Peter Schwabe, and Zooko Wilcox-O'Hearn. 2015. SPHINCS: Practical Stateless Hash-Based Signatures. In *Advances in Cryptology – EUROCRYPT (Lecture Notes in Computer Science)*, Vol. 9056. Springer, 368–397.
[6] Kemal Bicakci, Gene Tsudik, and Brian Tung. 2003. How to construct optimal one-time signatures. *Journal of Computer Networks* 43, 3 (2003), 339–349.
[7] Johannes Buchmann, Erik Dahmen, and Andreas Hülsing. 2011. XMSS - A Practical Forward Secure Signature Scheme Based on Minimal Security Assumptions. In *Post-Quantum Cryptography – PQCrypto (Lecture Notes in Computer Science)*, Vol. 7071. Springer, 117–129.
[8] Johannes Buchmann, Erik Dahmen, and Michael Schneider. 2008. Merkle Tree Traversal Revisited. In *PQCrypto (Lecture Notes in Computer Science)*, Vol. 5299. Springer, 63–78.
[9] Nicolas Courtois1, Matthieu Finiasz, and Nicolas Sendrier. 2001. How to Achieve a McEliece-based Digital Signature Scheme. In *Advances in Cryptology – ASIACRYPT (Lecture Notes in Computer Science)*, Vol. 2248. Springer, 157–174.
[10] Thomas M. Cover. 1973. Enumerative Source Encoding. *IEEE Transactions on Information Theory* 19, 1 (1973), 73–77.
[11] Erik Dahmen, Katsuyuki Okeya, Tsuyoshi Takagi, and Camille Vuillaume. 2008. Digital Signatures Out of Second-Preimage Resistant Hash Functions. In *PQCrypto (Lecture Notes in Computer Science)*, Vol. 5299. Springer, 109–123.
[12] Abhijit Das and C. E. Veni Madhavan. 2009. *Public-Key Cryptography: Theory and Practice.* Pearson.
[13] Ana Karin D. S. de Oliveira, Julio LÂ'opez, and Roberto Cabral. 2017. High Performance of Hash-based Signature Schemes. *International Journal of Advanced Computer Science and Applications* 8, 3 (2017).
[14] Whitfield Diffie and Martin E. Hellman. 1976. New directions in cryptography. *IEEE Transactions on Information Theory* 22, 6 (1976), 644–654.
[15] Léo Ducas, Alain Durmus, Tancrède Lepoint, and Vadim Lyubashevsky. 2013. Lattice Signatures and Bimodal Gaussians. In *Advances in Cryptology – CRYPTO (Lecture Notes in Computer Science)*, Vol. 8042. Springer, 40–56.
[16] Agner Fog. 2016. Software optimization resources. http://agner.org/optimize/.
[17] Andreas Hülsing. 2019. XMSS and XMSS^{MT}. https://huelsing.net/wordpress/?page_id=54.
[18] Andreas Hülsing. 2019. xmss-reference. https://github.com/joostrijneveld/xmss-reference.
[19] A. Hülsing, D. Butin, S.-L. Gazdag, and A. Mohaisen. 2017. XMSS: Extended Hash-Based Signatures. draft-irtf-cfrg-xmss-hash-based-signatures, work in progress.
[20] Andreas Hülsing, L. Rausch, and J. Buchmann. 2013. Optimal Parameters for XMSS^{MT}. In *Security Engineering and Intelligence Informatics: CD-ARES (Lecture Notes in Computer Science)*, Vol. 8128. Springer, 194–208.
[21] Andreas Hülsing, Joost Rijneveld, and Fang Song. 2016. Mitigating Multi-target Attacks in Hash-Based Signatures. In *Public Key Cryptography – PKC (Lecture Notes in Computer Science)*, Vol. 9614. Springer, 387–416.
[22] Intel. 2019. https://software.intel.com/sites/landingpage/IntrinsicsGuide/#.
[23] Kassem Kalach and Reihaneh Safavi-Naini. 2016. An Efficient Post-Quantum One-Time Signature Scheme. In *Selected Areas in Cryptography – SAC (LNCS)*, Vol. 9566. Springer, 331–351.
[24] Sabyasachi Karati. 2019. K2SN-MSS. https://github.com/skarati/K2SN-MSS.
[25] Sabyasachi Karati and Reihaneh Safavi-Naini. 2019. K2SN-MSS: An Efficient Post-Quantum Signature (Full Version). Cryptology ePrint Archive, Report 2019/442.
[26] VAMPIRE lab. 2019. Supercop: Version 2017.12.18. https://bench.cr.yp.to/supercop.html.
[27] Leslie Lamport. 1979. Constructing Digital Signatures from a One Way Function. technical report of SRI International.
[28] Patrick Longa and Michael Naehrig. 2016. Speeding up the Number Theoretic Transform for Faster Ideal Lattice-Based Cryptography. In *Cryptology and Network Security – CANS (Lecture Notes in Computer Science)*, Vol. 10052. Springer, 124–139.
[29] Vadim Lyubashevsky and Daniele Micciancio. 2006. Generalized Compact Knapsacks Are Collision Resistant. In *International Colloquium on Automata, Languages, and Programming – ICALP (Lecture Notes in Computer Science)*, Vol. 4052. Springer, 144–155.
[30] Vadim Lyubashevsky, Daniele Micciancio, Chris Peikert, and Alon Rosen. 2008. SWIFFT: A Modest Proposal for FFT Hashing. In *Fast Software Encryption – FSE (Lecture Notes in Computer Science)*, Vol. 5086. Springer, 54–72.
[31] Ralph C. Merkle. 1987. A Digital Signature Based on a Conventional Encryption Function. In *Advances in Cryptology – CRYPTO (Lecture Notes in Computer Science)*, Vol. 293. Springer, 369–378.
[32] Ralph C. Merkle. 1989. A Certified Digital Signature. In *Advances in Cryptology – CRYPTO (Lecture Notes in Computer Science)*, Vol. 435. Springer, 218–238.
[33] Daniele Micciancio. 2007. Generalized Compact Knapsacks, Cyclic Lattices, and Efficient One-Way Functions. *Computational Complexity* 16, 4 (2007), 365–411.
[34] Daniele Micciancio. 2008. https://github.com/micciancio/SWIFFT.
[35] Dustin Moody. 2018. Stateful hash-based signatures. https://groups.google.com/a/list.nist.gov/forum/#!topic/pqc-forum/WN3MJoe1RKQ.
[36] NIST. 2017. Post-Quantum Cryptography - Call for Proposals. https://csrc.nist.gov/Projects/Post-Quantum-Cryptography/Post-Quantum-Cryptography-Standardization/Call-for-Proposals.
[37] Chris Peikert and Alon Rosen. 2006. Efficient Collision-Resistant Hashing from Worst-Case Assumptions on Cyclic Lattices. In *Theory of Cryptography Conference – TCC (Lecture Notes in Computer Science)*, Vol. 3876. Springer, 145–166.
[38] John Rompel. 1990. One-way functions are necessary and sufficient for secure signatures. In *ACM symposium on Theory of computing – STOC*. ACM, 387–394.
[39] Peter W. Shor. 1994. Algorithms for quantum computation: discrete logarithms and factoring. In *Proceedings 35th Annual Symposium on Foundations of Computer Science*. 124–134.
[40] Tom Simonite. 2016. MIT Technology Review. https://www.technologyreview.com/s/600715/nsa-says-it-must-act-now-against-the-quantum-computing-threat/.
[41] Wanzhong Sun, Hongpeng Guo, Huilei He, and Zibin Dai. 2007. Design and optimized implementation of the SHA-2(256, 384, 512) hash algorithms. In *International Conference on ASIC*. IEEE, 858–861.
[42] Gregory M. Zaverucha and Douglas R. Stinson. 2011. Short one-time signatures. *Advances in Mathematics of Communications–AMC* 5, 3 (2011), 473–488.

A GENERALIZED SWIFFT-2^δ

This section includes the algorithm of gSWIFFT computation.

Algorithm 9 generalized SWIFFT-2^δ, in short gSWIFFT-2^δ

Input: $(\mathbf{z}_0, \ldots, \mathbf{z}_{15})$ where $\mathbf{z}_i \in \mathbb{Z}_2^{64}$ and $\left(\mathbf{A}_{i,0}, \ldots, \mathbf{A}_{i,\frac{64}{2^\delta}}\right)$ for $0 \leqslant i \leqslant 15$;

Output: $\mathbf{x}_S = \left(\mathbf{x}_{S,0}, \ldots, \mathbf{x}_{S,\frac{64}{2^\delta}-1}\right) \in \mathbb{Z}_{257}^{64}$;

---

(1) Compute $\mathbf{y}_i = g\mathrm{NTT2}^\delta(\mathbf{z}_i) = \left(\mathbf{y}_{i,0}, \ldots, \mathbf{y}_{i,\frac{64}{2^\delta}-1}\right), \forall 0 \leqslant i \leqslant 15$.

(2) Compute $\mathbf{t}_{i,j} = \mathbf{y}_{i,j} \odot_{257} \mathbf{A}_{i,1}, \forall 0 \leqslant i \leqslant 15, \forall 0 \leqslant j \leqslant \frac{64}{2^\delta}-1$.

(3) Compute $\mathbf{x}_{S,i} = \mathbf{x}_{S,i} \oplus_{257} \mathbf{t}_{j,i}, \forall 0 \leqslant i \leqslant \frac{64}{2^\delta}-1, \forall 0 \leqslant j \leqslant 15$.

B   PROOF SKETCH OF THEOREM 3.5

Due to the lack of spaces, we could not able to include the proof. Here we only include the over view of the security game and we refer the full version [25] for the detailed proof. Let $\mathcal{B}$ be a attacker who has $q_1 + q_2$ pairs of (x_i, y_i, k_i) pairs where $y_i = H_{k_i}(x_i)$ of the hash function family SWIFFT and must find a second preimage of one of them where $q_1 \leqslant 2^h$ and $q_2 < 2^{h+\ell}$. Let $\mathcal{A}$ be a SUF-CMA attacker against the K2SN-MSS signature scheme and has success probability ϵ to forge a signature after q queries in time τ. In the security game, $\mathcal{B}$ uses $\mathcal{A}$ as a sub-routine to find a second preimage of a tuple of the target list.

The game has three steps: (i) setup phase, (ii) query phase, and (iii) answer phase. In the setup phase, $\mathcal{B}$ creates an instance of K2SN-MSS signature scheme to sign 2^h messages, where each KSN-OTS has 2^ℓ secret key (and public-key) components. $\mathcal{B}$ will embed all the target pairs in the signature tree as follows: first randomly chooses q_1 KSN-OTSs and embeds q_1 target pairs in those KSN-OTSs. It then chooses q_2 positions in the Merkle tree of height $h + \ell$, uniformly at random, and embeds the remaining target pairs at those nodes.

Embedding a target pair (x', y', k') in KSN-OTS. First set the hash key of the KSN-OTS at k'. Randomly select 2^ℓ component secret keys and compute the respective component public keys. Now choose one of the 2^ℓ components uniformly at random and set the particular component secret key and corresponding component public key as x' and y' respectively.

Embedding a target pair (x', y', k') in Merkle Tree. Let the node be $y_{i,j}$. We set the hash key of the node $k_{i,j}$ at k' and the hash value $y_{i,j}$ at y'. Now we compute the random pad $\mathbf{v}_{i,j}$ as

$$\mathbf{v}_{i,j} = x' \oplus (\text{Merge}(y_{i+1,2j}, y_{i+1,2j+1})).$$

During the construction of the Merkle tree, $\mathcal{B}$ uses Merge instead of concatenation. Lemma 3.4 shows that Merge function is indistinguishable from concatenation function if the hash functions are modeled as random oracles. At the end $\mathcal{B}$ publishes the root node of the tree and all the hash keys and the random pads used in the tree. Notice that there are 2^h instances of KSN-OTSs and the Merkle tree of height $(h + \ell)$ has $2^{h+\ell} - 1$ intermediate nodes including the root nodes. Therefore, the upper bounds on q_1 and q_2 guarantee that successful embedding can be achieved.

Next phase is the query phase. In this phase, $\mathcal{A}$ asks signature queries for messages $M_i, 1 = 1 \cdots q$, to $\mathcal{B}$. $\mathcal{A}$ can query at most one query for each KSN-OTS instance, and therefore $\mathcal{A}$ is allowed to have at most 2^h queries. Let $\{(M_0, \sigma_0), (M_1, \sigma_1), \ldots, (M_{2^h-1}, \sigma_{2^h-1})\}$ denote the set of received message and signature pairs by $\mathcal{A}$.

After receiving responses to all the queries, $\mathcal{A}$ returns a forged message and signature pair, (M^*, σ^*). $\mathcal{A}$ is successful if $(M^*, \sigma^*) \neq (M_i, \sigma_i)$ and K2SN.verify$(M^*, \sigma^*) = 1$ for all $0 \leqslant i < 2^h$. By assumption, this happens with probability ϵ. We will show that $\mathcal{B}$ can use this forgery to find a second pre-image for the hash function with success probability given by ϵ'. We distinguish between the following cases: (i) the forged signature pair (M^*, σ^*) is in one of the q_1 chosen KSN-OTSs by $\mathcal{B}$, and (ii) the forged signature (M^*, σ^*) is in one of the remaining KSN-OTSs. Let the index of the KSN-OTS (received as the part of the signature σ^*) be i^*. Each of the above two case can be further subdivide into three cases: (i) $M^* = M_{i^*}$

and $\sigma^* \neq \sigma_{i^*}$, (ii)$M^* \neq M_{i^*}$ and $\sigma^* = \sigma_{i^*}$, and (iii)$M^* \neq M_{i^*}$ and $\sigma^* \neq \sigma_{i^*}$. Notice that these cases are mutually exclusive and so the maximum success probability of $\mathcal{B}$ is the maximum probability of these cases. We show that

$$\epsilon' = (1 - 2^{-2n} - 2^{-m})(1 - 2^{-2n}) \cdot \max\left\{2^{-2h}2^{-2(h+\ell)}, (1 - 2^{-2h})2^{-2(h+\ell)}\right\}.$$

Details are omitted because of space. Therefore, $\mathcal{B}$ can find a second preimage with probability $\epsilon\epsilon'$, and upper bound of this InSec$^{\text{MM-SPR}}$ $(\mathcal{H}, 2^h, 2^h(1 + 2^\ell))$. Therefore, the upper bound of ϵ becomes

$$\epsilon \leqslant \text{InSec}^{\text{MM-SPR}}(\mathcal{H}, 2^h, 2^h(1 + 2^\ell))/\epsilon'. \quad \square$$

C   OPTIMIZATION TABLES

Table 5: Final computations of $\mathbf{y}_{i,j_1}$

j_1	$\mathbf{y}_{i,j_1} \in \mathbb{Z}_{257}^8$	j_1	$\mathbf{y}_{i,j_1} \in \mathbb{Z}_{257}^8$
0	$\mathbf{y}''_{i,0} \oplus_{257} \mathbf{y}''_{i,1}$	4	$\mathbf{y}''_{i,0} \ominus_{257} \mathbf{y}''_{i,1}$
1	$\mathbf{y}''_{i,8} \oplus_{257} \mathbf{y}''_{i,9}$	5	$\mathbf{y}''_{i,8} \ominus_{257} \mathbf{y}''_{i,9}$
2	$\mathbf{y}''_{i,6} \oplus_{257} \mathbf{y}''_{i,7}$	6	$\mathbf{y}''_{i,6} \ominus_{257} \mathbf{y}''_{i,7}$
3	$\mathbf{y}''_{i,10} \oplus_{257} \mathbf{y}''_{i,11}$	7	$\mathbf{y}''_{i,10} \ominus_{257} \mathbf{y}''_{i,11}$

Table 6: Sub-computations of $\mathbf{y}_{i,j_1}$ of Table 5, where each $\mathbf{y}''_{i,.} \in \mathbb{Z}_{257}^8$. Component-wise left Shift is denoted by $\ll$.

$\mathbf{y}''_{i,0} = \mathbf{y}'_{i,0} \oplus_{257} \mathbf{y}'_{i,2} \oplus_{257} \mathbf{y}'_{i,4} \oplus_{257} \mathbf{y}'_{i,6}$
$\mathbf{y}''_{i,1} = \mathbf{y}'_{i,1} \oplus_{257} \mathbf{y}'_{i,3} \oplus_{257} \mathbf{y}'_{i,5} \oplus_{257} \mathbf{y}'_{i,7}$
$\mathbf{y}''_{i,2} = \mathbf{y}'_{i,0} \ominus_{257} \mathbf{y}'_{i,4}$
$\mathbf{y}''_{i,3} = \left(\mathbf{y}'_{i,1} \ominus_{257} \mathbf{y}'_{i,5}\right)$
$\mathbf{y}''_{i,4} = \left(\mathbf{y}'_{i,2} \ominus_{257} \mathbf{y}'_{i,6}\right) \ll 4$
$\mathbf{y}''_{i,5} = \left(\mathbf{y}'_{i,3} \ominus_{257} \mathbf{y}'_{i,7}\right)$
$\mathbf{y}''_{i,6} = \mathbf{y}'_{i,0} \oplus_{257} \mathbf{y}'_{i,4} \ominus_{257} \mathbf{y}'_{i,2} \ominus_{257} \mathbf{y}'_{i,6}$
$\mathbf{y}''_{i,7} = \left(\mathbf{y}'_{i,1} \oplus_{257} \mathbf{y}'_{i,5} \ominus_{257} \mathbf{y}'_{i,3} \ominus_{257} \mathbf{y}'_{i,7}\right) \ll 4$
$\mathbf{y}''_{i,8} = \mathbf{y}''_{i,2} \oplus_{257} \mathbf{y}''_{i,4}$
$\mathbf{y}''_{i,9} = (\mathbf{y}''_{i,3} \ll 2) \oplus_{257} (\mathbf{y}''_{i,5} \ll 6)$
$\mathbf{y}''_{i,10} = \mathbf{y}''_{i,2} \ominus_{257} \mathbf{y}''_{i,4}$
$\mathbf{y}''_{i,11} = (\mathbf{y}''_{i,3} \ll 6) \oplus_{257} (\mathbf{y}''_{i,5} \ll 2)$

Table 7: Final computations of $\mathbf{y}_{i,j_1}$

j_1	$\mathbf{y}_{i,j_1} \in \mathbb{Z}_{257}^{16}$
0	$\mathbf{y}''_{i,0} \oplus_{257} \mathbf{y}''_{i,1}$
1	$\mathbf{y}''_{i,2} \oplus_{257} \mathbf{y}''_{i,3}$
2	$\mathbf{y}''_{i,0} \ominus_{257} \mathbf{y}''_{i,1}$
3	$\mathbf{y}''_{i,2} \ominus_{257} \mathbf{y}''_{i,3}$

Table 8: Sub-computations of $\mathbf{y}_{i,j_1}$ of Table 7, where each $\mathbf{y}''_{i,.} \in \mathbb{Z}_{257}^{16}$. Component-wise left Shift is denoted by $\ll$.

$\mathbf{y}''_{i,0} = \mathbf{y}'_{i,0} \oplus_{257} \mathbf{y}'_{i,2}$
$\mathbf{y}''_{i,1} = \mathbf{y}'_{i,1} \oplus_{257} \mathbf{y}'_{i,3}$
$\mathbf{y}''_{i,2} = \mathbf{y}'_{i,0} \ominus_{257} \mathbf{y}'_{i,2}$
$\mathbf{y}''_{i,3} = (\mathbf{y}'_{i,1} \ominus_{257} \mathbf{y}'_{i,3}) \ll 4$

D SWIFFT-16-AVX2 IMPLEMENTATION

D.1 Modular Reduction

We use two types of modular reductions referred to as LazyReduced16 and Reduced16. In LazyReduced16 the output vector can have components that are negative integers, while the components of Reduced16 output vector are integers in the range $[0, p - 1]$. Even though LazyReduced16 is a faster operation, the negative output may not be accepted as the input in the next operation such as shifting. Thus we use LazyReduced16 when negative output is acceptable by the next operation.

(1) LazyReduced16: Let $\mathbf{c} = (c_0, c_1, \ldots, c_{15})$. Each component c_i of $\mathbf{c}$ can be written as $c_i = c_{i,0} + c_{i,1} \cdot 2^8 = (c_{i,0} - c_{i,1}) + c_{i,1} \cdot (2^8 + 1)$. Therefore $c_i = c_{i,0} - c_{i,1} \pmod{p}$. Note that the i-th component, depending on the values of $c_{i,0}$ and $c_{i,1}$, can be negative or positive. The _mm256_and_si256 instruction computes the bit-wise AND of two __m256i registers. Let $mask255$ be a vector of length 16, and each component be 255. Applying _mm256_and_si256 on $\mathbf{c}$ and $mask255$, we obtain the vector $\mathbf{c}_0 = (c_{0,0}, c_{1,0}, \ldots, c_{15,0})$. The _mm256_srli_epi16 performs right-shift on packed 16-bit integers of __m256i. We compute $\mathbf{c}_1 = $ _mm256_srli _epi16($\mathbf{c}, 8$). Let $\mathbf{c}_1 = (c_{0,1}, c_{1,1}, \ldots, c_{15,1})$. We obtain the lazy reduced vector $\hat{\mathbf{c}}$ as _mm256_sub_epi16($\mathbf{c}_0, \mathbf{c}_1$), where _mm256_sub_epi16 performs subtraction on packed 16-bit integers of two __m256i registers.

(2) Reduced16: Let $\hat{\mathbf{c}} = \{\hat{c}_0, \hat{c}_1, \ldots, \hat{c}_{15}\} = $ LazyReduced16($\mathbf{c}$). To determine if the i-th component of the reduced vector $\hat{\mathbf{c}}$ is negative, we use $\mathbf{c}' = $ _mm256_cmpgt_epi16($\hat{\mathbf{c}}$, allone), where allone is a vector of all -1 component. If $\hat{c}_i \geqslant 0$, then $\mathbf{c}'_i = -1$, else it is 0. Thus we obtain $\mathbf{c}^*$ with 0 and -1 components, and 0 components corresponding to negative components of $\mathbf{c}'$. Let $\mathbf{p}$ be a vector of length 16, each component being $p = 257$. By applying _mm256_and_si256 on $\mathbf{p}$ and $\mathbf{c}'$, we obtain a vector whose i-th component is p if $\hat{c}_i < 0$, else 0. To obtain $\hat{\mathbf{c}}$ modulo p, we use _mm256_add_epi16($\mathbf{p}'$ and $\hat{\mathbf{c}}$), to add $\mathbf{p}'$ and $\hat{\mathbf{c}}$ where _mm256_add_epi16 performs additions on packed 16-bit integers of the two __m256i registers.

For two vectors $\mathbf{a} = (a_0, a_1, \ldots, a_{15})$ and $\mathbf{b} = (b_0, b_1, \ldots, b_{15})$, we perform the following vector operations.

D.2 Modular Vector Addition ($\oplus_{257}$)

The addition of the vectors is performed by instruction _mm256_add _epi16and then perform the modular reduction on the result.

D.3 Modular Vector Subtraction ($\ominus_{257}$)

First add vector $k\mathbf{p}$ (each component kp, for an appropriate choice of integer k), to vector $\mathbf{a}$ such that all components of the vector $(\mathbf{a} + k\mathbf{p}) - \mathbf{b}$ are positive. The vector subtraction operation is done by the instruction _mm256_sub_epi16. Because of adding $k\mathbf{p}$, the positive components of $(\mathbf{a} - \mathbf{b})$ will become large and so we perform a modular reduction to reduce each component to 9-bit.

D.4 Modular Vector Multiplication ($\odot_{257}$)

Each a_i and b_i are 9-bit long and so $c_i = a_i \cdot b_i$ may need more than 16-bit to represent, and this will result in an overflow. To overcome this problem, c_i is divided into two parts, $c_i = c_{i,0} + c_{i,1} \cdot 2^{16}$, and can be further reduced as $c_i = c_{i,0} + c_{i,1} \pmod{257}$.

The maximum value of a_i and b_i is 256. If both a_i and b_i are 256 then only $c_{i,1}$ becomes 1 and $c_{i,0}$ is zero. For all other combinations of a_i and b_i, $c_{i,1}$ is 0.

Let $\mathbf{c}_0 = \{c_{0,0}, c_{1,0}, \ldots, c_{15,0}\}$ and $\mathbf{c}_1 = \{c_{0,1}, c_{1,1}, \ldots, c_{15,1}\}$. We compute vector $\mathbf{c}_0$ using the instruction _mm256_mullo_epi16, and $\mathbf{c}_1$ using the instruction _mm256_mulhi_epu16 from $\mathbf{a}$ and $\mathbf{b}$. Then we perform modular reduction Reduced16 on the vector $\mathbf{c}_0$. The final vector is $\mathbf{c}^*$. Adding $\mathbf{c}^*$ and $\mathbf{c}_1$ gives the desired reduced vector.

D.5 Implementation of the function Select16

The Select function takes an input vector of length 16 and outputs a vector of length 16 whose each component is the same as the one of the selected component of the input vector by an index as described in Section 4.2. This has been implemented using _mm256_permutevar8x32_epi32 as avx2 does not provide a permutation operation on 16-bit __m256i register data. So, implementations of Select16() and Select8() are different. For Select16(), we use the set of vectors $\{\mathbf{s}_0, \mathbf{s}_1, \ldots, \mathbf{s}_7\}$ where each $\mathbf{s}_i, i = 0, 1, \ldots, 7$ is an 8-dimensional vector with each component i as $\mathbf{s}_i = (i, i, \ldots, i)$. Let $\mathbf{a} = (a_{2\cdot0}, a_{2\cdot0+1}, a_{2\cdot1}, a_{2\cdot1+1}, \ldots, a_{2\cdot7}, a_{2\cdot7+1})$ be the input to Select() function. First we select the i-th pair $(a_{2\cdot i}, a_{2\cdot i+1})$ of $\mathbf{a}$ using the instruction _mm256_permutevar8x32_epi32, and construct the vector $\mathbf{a}_i$ as $\mathbf{a}_i = $ _mm256_permutevar8x32_epi32($\mathbf{a}_i, \mathbf{s}_i$) $= (a_{2\cdot i}, a_{2\cdot i+1}, \ldots)$.

To select the $2i$-th component of $\mathbf{a}$, we first compute the vector $\mathbf{a}_{i,0} = (a_{2\cdot i}, 0, \ldots, a_{2\cdot i}, 0)$ using the instruction _mm256_and_si256 on vectors $\mathbf{a}_i$ and vector $(-1, 0, \ldots, -1, 0)$. Applying _mm256_srli _epi32 on $\mathbf{a}_{i,0}$ for 16-bit, we get $\mathbf{a}'_{i,0} = (0, a_{2\cdot i}, \ldots, 0, a_{2\cdot i})$. If we perform _mm256_or_si256 on $\mathbf{a}_{i,0}$ and $\mathbf{a}'_{i,0}$, the resulting vector will be the output of the select function for $2i$. Similarly, we perform the select operation for $(2i + 1)$-th component.

D.6 Further Details

Reduction is a costly operation. In the following, we show how the number of reductions can be reduced. In bNTT16 computation of Algorithm 8, multiplications in step 3 are accompanied by a LazyReduced16 and so the components of the vectors $\mathbf{y}'_{i,j}$ are in the range $[-256, +255]$. To compute the vectors $\mathbf{y}_{i,j}$ in step 4, we do not perform reduction for vector addition and subtraction, and subtraction uses $k\mathbf{p}$ for $k = 2^2$ and 2^5. This results in the components of the output vector to be at most 14-bit long. The multiplications $\mathbf{t}'_{i,j} = \mathbf{y}_{i,j} \odot_{257} \mathbf{A}'_{i,j}$ in Algorithm 8 are implemented using the above vector multiplication followed by a Reduced16 operation, resulting in each component of the output vector to be 9-bit and so no further reduction for the additions $\mathbf{x}'_{S,i} = \mathbf{x}'_{S,i} \oplus_{257} \mathbf{t}'_{i,j}$ is needed. This results in vectors whose components are at most 13 bits long. Finally, we use Reduced16 on the vector $\mathbf{x}'_{S,i}$ to produce the final result where components of $\mathbf{x}'_{S,i}$ are all at most 9 bits.

These optimizations reduce the number of reduction operations. Tables 5–8 show that the total number of required vector operations for steps 3 and 4 of 16-way parallelization of NTT which are significantly smaller than the number of operations required for 8-way parallelization, and this leads to a substantial speedup of the implementation of SWIFFT-16-avx2.